

Applications of Machine Learning - 101

Alfonso R. Reyes

2019-05-15

Contents

1	Prerequisites	5
2	PCA: prcomp vs princomp	7
2.1	General methods for principal component analysis	7
2.2	prcomp() and princomp() functions	7
2.3	factoextra	7
2.4	demo dataset	8
2.5	Compute PCA in R using prcomp()	8
2.6	Plots: quality and contribution	9
2.7	Access to the PCA results	12
3	Predict using PCA	19
3.1	Supplementary variables	21
3.2	Theory behind PCA results	23
4	Diagnostic Plots	27
4.1	Residual vs Fitted plot	29
4.2	Normal QQ	29
4.3	Scale-location	32
4.4	Cook's Distance	33
4.5	Residual vs Leverage Plot	34
4.6	Cook's dist vs Leverage $h_{ii}/(1-h_{ii})$	35
5	Temperature modeling using nested dataframes	37
5.1	Prepare the data	37
5.2	Define the models	41
5.3	Test modeling on one dataset	42
5.4	Making a nested dataframe	45
5.5	Apply multiple models on a nested structure	47
5.6	Using broom package to look at model-statistics	53
6	Logistic Regression. Diabetes	55
6.1	Introduction	55
6.2	Exploring the data	55
6.3	Logistic regression with R	57
6.4	A second model	59
6.5	Classification rate and confusion matrix	60
6.6	Plots and decision boundaries	61

Chapter 1

Prerequisites

This is a *sample* book written in **Markdown**. You can use anything that Pandoc's Markdown supports, e.g., a math equation $a^2 + b^2 = c^2$.

The **bookdown** package can be installed from CRAN or Github:

```
install.packages("bookdown")  
# or the development version  
# devtools::install_github("rstudio/bookdown")
```

Remember each Rmd file contains one and only one chapter, and a chapter is defined by the first-level heading #.

To compile this example to PDF, you need XeLaTeX. You are recommended to install TinyTeX (which includes XeLaTeX): <https://yihui.name/tinytex/>.

Chapter 2

PCA: prcomp vs princomp

<http://www.sthda.com/english/articles/31-principal-component-methods-in-r-practical-guide/118-principal-component-analysis>

2.1 General methods for principal component analysis

There are two general methods to perform PCA in R :

- Spectral decomposition which examines the covariances / correlations between variables
- Singular value decomposition which examines the covariances / correlations between individuals

The function `princomp()` uses the spectral decomposition approach. The functions `prcomp()` and `PCA()`[FactoMineR] use the singular value decomposition (SVD).

2.2 prcomp() and princomp() functions

The simplified format of these 2 functions are :

```
prcomp(x, scale = FALSE)
princomp(x, cor = FALSE, scores = TRUE)
```

1. Arguments for `prcomp()`:
x: a numeric matrix or data frame
scale: a logical value indicating whether the variables should be scaled to have unit variance before the analysis takes place
2. Arguments for `princomp()`:
x: a numeric matrix or data frame **cor**: a logical value. If TRUE, the data will be centered and scaled before the analysis **scores**: a logical value. If TRUE, the coordinates on each principal component are calculated

2.3 factoextra

```
# install.packages("factoextra")
```

```
library(factoextra)
```

2.4 demo dataset

We'll use the data sets `decathlon2` [in `factoextra`], which has been already described at: PCA - Data format.

Briefly, it contains:

- Active individuals (rows 1 to 23) and active variables (columns 1 to 10), which are used to perform the principal component analysis
- Supplementary individuals (rows 24 to 27) and supplementary variables (columns 11 to 13), which coordinates will be predicted using the PCA information and parameters obtained with active individuals/variables.

```
library("factoextra")
data(decathlon2)
decathlon2.active <- decathlon2[1:23, 1:10]
head(decathlon2.active[, 1:6])
```

#>		X100m	Long.jump	Shot.put	High.jump	X400m	X110m.hurdle
#>	SEBRLE	11.04	7.58	14.83	2.07	49.81	14.69
#>	CLAY	10.76	7.40	14.26	1.86	49.37	14.05
#>	BERNARD	11.02	7.23	14.25	1.92	48.93	14.99
#>	YURKOV	11.34	7.09	15.19	2.10	50.42	15.31
#>	ZSIVOCZKY	11.13	7.30	13.48	2.01	48.62	14.17
#>	McMULLEN	10.83	7.31	13.76	2.13	49.91	14.38

```
decathlon2.supplementary <- decathlon2[24:27, 1:10]
head(decathlon2.supplementary[, 1:6])
```

#>		X100m	Long.jump	Shot.put	High.jump	X400m	X110m.hurdle
#>	KARPOV	11.02	7.30	14.77	2.04	48.37	14.09
#>	WARNERS	11.11	7.60	14.31	1.98	48.68	14.23
#>	Nool	10.80	7.53	14.26	1.88	48.81	14.80
#>	Drews	10.87	7.38	13.07	1.88	48.51	14.01

2.5 Compute PCA in R using prcomp()

In this section we'll provide an easy-to-use R code to compute and visualize PCA in R using the `prcomp()` function and the `factoextra` package.

1. Load `factoextra` for visualization

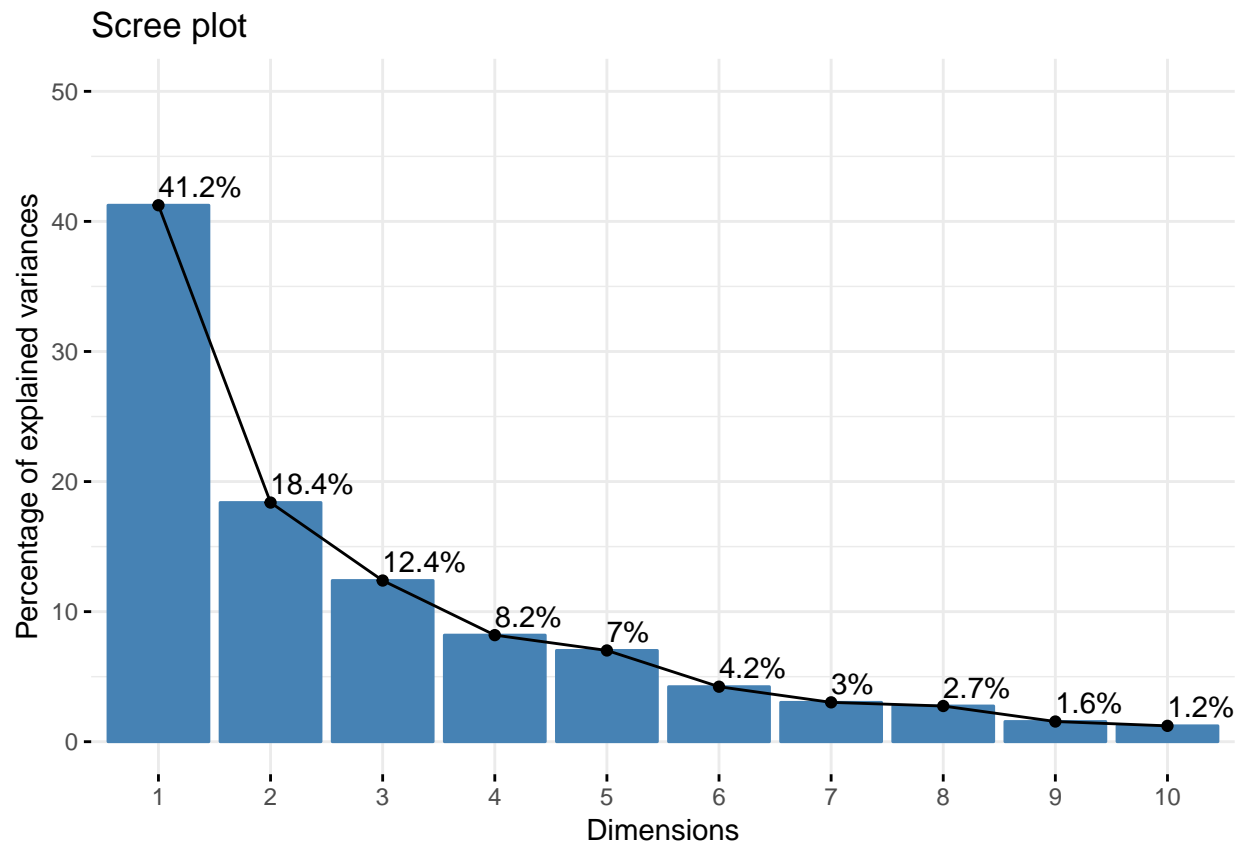
```
library(factoextra)
```

2. compute PCA

```
# compute PCA
res.pca <- prcomp(decathlon2.active, scale = TRUE)
```

3. Visualize eigenvalues (scree plot). Show the percentage of variances explained by each principal component.

```
# Visualize eigenvalues (scree plot).
fviz_eig(res.pca, addlabels = TRUE, ylim = c(0, 50))
```

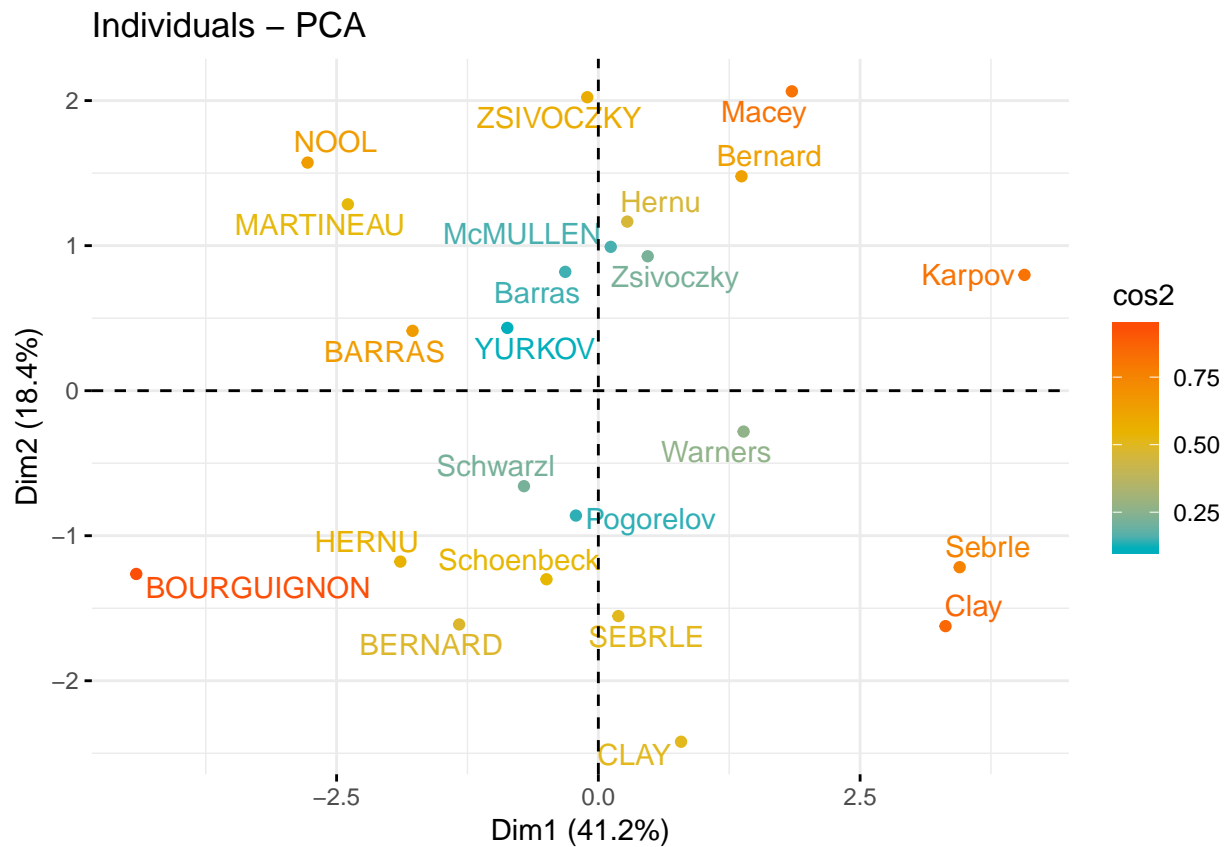



From the plot above, we might want to stop at the fifth principal component. 87% of the information (variances) contained in the data are retained by the first five principal components.

2.6 Plots: quality and contribution

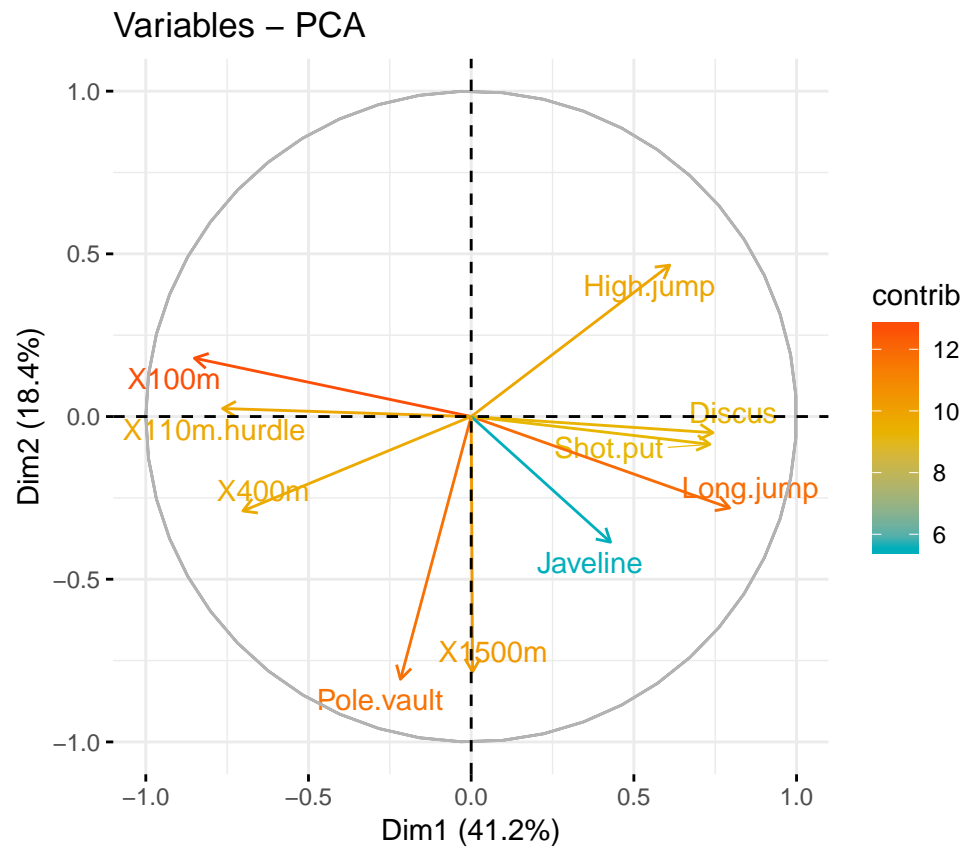
4. Graph of individuals. Individuals with a similar profile are grouped together.

```
# Graph of individuals.
fviz_pca_ind(res.pca,
  col.ind = "cos2", # Color by the quality of representation
  gradient.cols = c("#00AFBB", "#E7B800", "#FC4E07"),
  repel = TRUE      # Avoid text overlapping
)
```



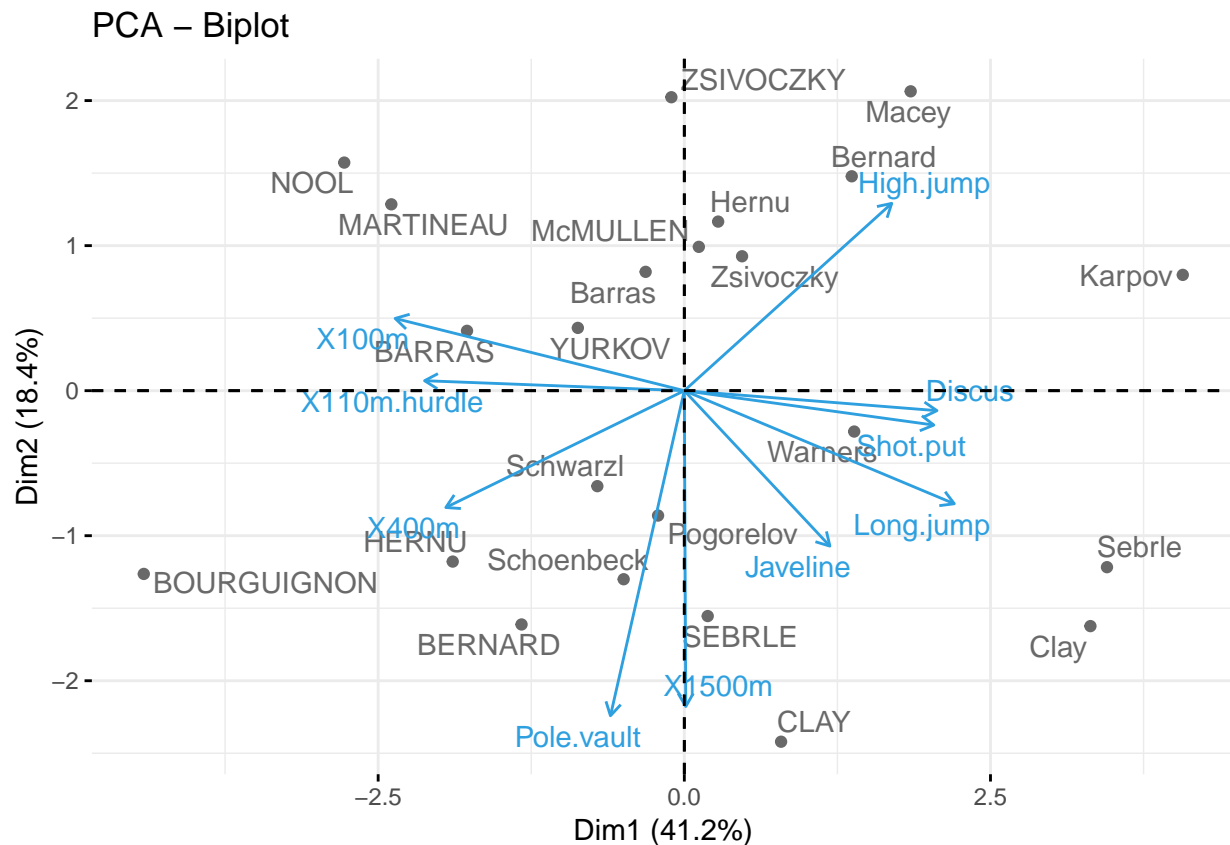
5. Graph of variables. Positive correlated variables point to the same side of the plot. Negative correlated variables point to opposite sides of the graph.

```
# Graph of variables.
fviz_pca_var(res.pca,
  col.var = "contrib", # Color by contributions to the PC
  gradient.cols = c("#00AFBB", "#E7B800", "#FC4E07"),
  repel = TRUE        # Avoid text overlapping
)
```



6. Biplot of individuals and variables

```
# Biplot of individuals and variables
fviz_pca_biplot(res.pca, repel = TRUE,
  col.var = "#2E9FDF", # Variables color
  col.ind = "#696969" # Individuals color
)
```



2.7 Access to the PCA results

```
library(factoextra)
# Eigenvalues
eig.val <- get_eigenvalue(res.pca)
eig.val
```

```
#>      eigenvalue variance.percent cumulative.variance.percent
#> Dim.1    4.1242133         41.242133          41.24213
#> Dim.2    1.8385309         18.385309          59.62744
#> Dim.3    1.2391403         12.391403          72.01885
#> Dim.4    0.8194402          8.194402          80.21325
#> Dim.5    0.7015528          7.015528          87.22878
#> Dim.6    0.4228828          4.228828          91.45760
#> Dim.7    0.3025817          3.025817          94.48342
#> Dim.8    0.2744700          2.744700          97.22812
#> Dim.9    0.1552169          1.552169          98.78029
#> Dim.10   0.1219710          1.219710         100.00000
```

```
# Results for Variables
```

```
res.var <- get_pca_var(res.pca)
res.var$coord # Coordinates
```

```
#>      Dim.1    Dim.2    Dim.3    Dim.4    Dim.5
#> X100m   -0.850625692  0.17939806 -0.30155643  0.03357320 -0.1944440
```

```

#> Long.jump      0.794180641 -0.28085695  0.19054653 -0.11538956  0.2331567
#> Shot.put       0.733912733 -0.08540412 -0.51759781  0.12846837 -0.2488129
#> High.jump      0.610083985  0.46521415 -0.33008517  0.14455012  0.4027002
#> X400m          -0.701603377 -0.29017826 -0.28353292  0.43082552  0.1039085
#> X110m.hurdle -0.764125197  0.02474081 -0.44888733 -0.01689589  0.2242200
#> Discus         0.743209016 -0.04966086 -0.17652518  0.39500915 -0.4082391
#> Pole.vault    -0.217268042 -0.80745110 -0.09405773 -0.33898477 -0.2216853
#> Javeline       0.428226639 -0.38610928 -0.60412432 -0.33173454  0.1978128
#> X1500m         0.004278487 -0.78448019  0.21947068  0.44800961  0.2632527
#>               Dim.6      Dim.7      Dim.8      Dim.9
#> X100m          0.035374780 -0.091336386 -0.104716925 -0.30306448
#> Long.jump     -0.033727883 -0.154330810 -0.397380703 -0.05158951
#> Shot.put      -0.239789034 -0.009886612  0.024359049  0.04778655
#> High.jump     -0.284644846  0.028157465  0.084405578 -0.11213822
#> X400m         -0.049289996  0.286106008 -0.233552216  0.08216041
#> X110m.hurdle  0.002632395 -0.370072158 -0.008344682  0.16176025
#> Discus        0.198544870 -0.142725641 -0.039559255  0.01336209
#> Pole.vault    -0.327464549 -0.010393176  0.032914942 -0.02576874
#> Javeline      0.362097598  0.133564318  0.052841099 -0.04045397
#> X1500m        0.042050151 -0.111367083  0.194469730 -0.10224014
#>               Dim.10
#> X100m          0.044417974
#> Long.jump      0.029719453
#> Shot.put       0.217451948
#> High.jump     -0.133566774
#> X400m         -0.034170673
#> X110m.hurdle -0.015629914
#> Discus        -0.172590426
#> Pole.vault    -0.137211339
#> Javeline      -0.003854347
#> X1500m        0.062834809

```

```
res.var$contrib      # Contributions to the PCs
```

```

#>               Dim.1      Dim.2      Dim.3      Dim.4      Dim.5
#> X100m          1.754429e+01  1.7505098  7.3386590  0.13755240  5.389252
#> Long.jump      1.529317e+01  4.2904162  2.9300944  1.62485936  7.748815
#> Shot.put       1.306014e+01  0.3967224  21.6204325  2.01407269  8.824401
#> High.jump      9.024811e+00  11.7715838  8.7928883  2.54987951  23.115504
#> X400m          1.193554e+01  4.5799296  6.4876363  22.65090599  1.539012
#> X110m.hurdle  1.415754e+01  0.0332933  16.2612611  0.03483735  7.166193
#> Discus         1.339309e+01  0.1341398  2.5147385  19.04132022  23.755756
#> Pole.vault     1.144592e+00  35.4618611  0.7139512  14.02307063  7.005084
#> Javeline       4.446377e+00  8.1086683  29.4531777  13.42963254  5.577615
#> X1500m         4.438531e-04  33.4728757  3.8871610  24.49386930  9.878367
#>               Dim.6      Dim.7      Dim.8      Dim.9      Dim.10
#> X100m          0.295915322  2.75705260  3.99520353  59.1740009  1.61756139
#> Long.jump      0.269003613  7.87159392  57.53322220  1.7146826  0.72414393
#> Shot.put       13.596858744  0.03230371  0.21618512  1.4712015  38.76768578
#> High.jump      19.159607001  0.26202607  2.59565787  8.1015517  14.62649091
#> X400m          0.574509906  27.05274658  19.87344405  4.3489667  0.95730504
#> X110m.hurdle  0.001638634  45.26163460  0.02537025  16.8579392  0.20028870
#> Discus         9.321746508  6.73226823  0.57016606  0.1150295  24.42174410
#> Pole.vault     25.357622290  0.03569883  0.39472201  0.4278065  15.43559151
#> Javeline       31.004964393  5.89573984  1.01729950  1.0543458  0.01217993

```

```
#> X1500m          0.418133591  4.09893563 13.77872941  6.7344755  3.23700871
```

```
res.var$cos2          # Quality of representation
```

```
#>          Dim.1      Dim.2      Dim.3      Dim.4      Dim.5
#> X100m      7.235641e-01  0.0321836641 0.090936280 0.0011271597 0.03780845
#> Long.jump  6.307229e-01  0.0788806285 0.036307981 0.0133147506 0.05436203
#> Shot.put   5.386279e-01  0.0072938636 0.267907488 0.0165041211 0.06190783
#> High.jump  3.722025e-01  0.2164242070 0.108956221 0.0208947375 0.16216747
#> X400m      4.922473e-01  0.0842034209 0.080390914 0.1856106269 0.01079698
#> X110m.hurdle 5.838873e-01  0.0006121077 0.201499837 0.0002854712 0.05027463
#> Discus     5.523596e-01  0.0024662013 0.031161138 0.1560322304 0.16665918
#> Pole.vault 4.720540e-02  0.6519772763 0.008846856 0.1149106765 0.04914437
#> Javeline   1.833781e-01  0.1490803723 0.364966189 0.1100478063 0.03912992
#> X1500m     1.830545e-05  0.6154091638 0.048167378 0.2007126089 0.06930197
#>          Dim.6      Dim.7      Dim.8      Dim.9
#> X100m      1.251375e-03  0.0083423353 1.096563e-02 0.0918480768
#> Long.jump  1.137570e-03  0.0238179990 1.579114e-01 0.0026614779
#> Shot.put   5.749878e-02  0.0000977451 5.933633e-04 0.0022835540
#> High.jump  8.102269e-02  0.0007928428 7.124302e-03 0.0125749811
#> X400m      2.429504e-03  0.0818566479 5.454664e-02 0.0067503333
#> X110m.hurdle 6.929502e-06  0.1369534023 6.963371e-05 0.0261663784
#> Discus     3.942007e-02  0.0203706085 1.564935e-03 0.0001785453
#> Pole.vault 1.072330e-01  0.0001080181 1.083393e-03 0.0006640282
#> Javeline   1.311147e-01  0.0178394271 2.792182e-03 0.0016365234
#> X1500m     1.768215e-03  0.0124026272 3.781848e-02 0.0104530472
#>          Dim.10
#> X100m      1.972956e-03
#> Long.jump  8.832459e-04
#> Shot.put   4.728535e-02
#> High.jump  1.784008e-02
#> X400m      1.167635e-03
#> X110m.hurdle 2.442942e-04
#> Discus     2.978746e-02
#> Pole.vault 1.882695e-02
#> Javeline   1.485599e-05
#> X1500m     3.948213e-03
```

```
# Results for individuals
```

```
res.ind <- get_pca_ind(res.pca)
```

```
res.ind$coord          # Coordinates
```

```
#>          Dim.1      Dim.2      Dim.3      Dim.4      Dim.5
#> SEBRLE      0.1912074 -1.5541282 -0.62836882 0.08205241 1.1426139415
#> CLAY        0.7901217 -2.4204156 1.35688701 1.26984296 -0.8068483724
#> BERNARD    -1.3292592 -1.6118687 -0.19614996 -1.92092203 0.0823428202
#> YURKOV     -0.8694134 0.4328779 -2.47398223 0.69723814 0.3988584116
#> ZSIVOCZKY  -0.1057450 2.0233632 1.30493117 -0.09929630 -0.1970241089
#> McMULLEN   0.1185550 0.9916237 0.84355824 1.31215266 1.5858708644
#> MARTINEAU  -2.3923532 1.2849234 -0.89816842 0.37309771 -2.2433515889
#> HERNU      -1.8910497 -1.1784614 -0.15641037 0.89130068 -0.1267412520
#> BARRAS     -1.7744575 0.4125321 0.65817750 0.22872866 -0.2338366980
#> NOOL       -2.7770058 1.5726757 0.60724821 -1.55548081 1.4241839810
#> BOURGUIGNON -4.4137335 -1.2635770 -0.01003734 0.66675478 0.4191518468
#> Sebrle     3.4514485 -1.2169193 -1.67816711 -0.80870696 -0.0250530746
#> Clay       3.3162243 -1.6232908 -0.61840443 -0.31679906 0.5691645854
```

```

#> Karpov      4.0703560  0.7983510  1.01501662  0.31336354 -0.7974259553
#> Macey       1.8484623  2.0638828 -0.97928455  0.58469073 -0.0002157834
#> Warners     1.3873514 -0.2819083  1.99969621 -1.01959817 -0.0405401497
#> Zsivoczky   0.4715533  0.9267436 -1.72815525 -0.18483138  0.4073029909
#> Hernu       0.2763118  1.1657260  0.17056375 -0.84869401 -0.6894795441
#> Bernard     1.3672590  1.4780354  0.83137913  0.74531557  0.8598016482
#> Schwarzl    -0.7102777 -0.6584251  1.04075176 -0.92717510 -0.2887568007
#> Pogorelov   -0.2143524 -0.8610557  0.29761010  1.35560294 -0.0150531057
#> Schoenbeck  -0.4953166 -1.3000530  0.10300360 -0.24927712 -0.6452257128
#> Barras      -0.3158867  0.8193681 -0.86169481 -0.58935985 -0.7797389436
#>
#>          Dim.6      Dim.7      Dim.8      Dim.9      Dim.10
#> SEBRLE      -0.46389755 -0.20796012  0.043460568 -0.659344137  0.03273238
#> CLAY        1.30420016 -0.21291866  0.617240611 -0.060125359 -0.31716015
#> BERNARD     -0.40062867 -0.40643754  0.703856040  0.170083313 -0.09908142
#> YURKOV      0.10286344 -0.32487448  0.114996135 -0.109524039 -0.11969720
#> ZSIVOCZKY   0.89554111  0.08825624 -0.202341299 -0.523103099 -0.34842265
#> McMULLEN    0.18657283  0.47828432  0.293089967 -0.105623196 -0.39317797
#> MARTINEAU   -0.45666350 -0.29975522 -0.291628488 -0.223417655 -0.61640509
#> HERNU       0.43623496 -0.56609980 -1.529404317  0.006184409  0.55368016
#> BARRAS      0.09026010  0.21594095  0.682583078 -0.669282042  0.53085420
#> NOOL        0.49716399 -0.53205687 -0.433385655 -0.115777808 -0.09622142
#> BOURGUIGNON -0.08200220 -0.59833739  0.563619921  0.525814030  0.05855882
#> Sebrle      -0.08279306  0.01016177 -0.030585843 -0.847210682  0.21970353
#> Clay        0.77715960  0.25750851 -0.580638301  0.409776590 -0.61601933
#> Karpov      -0.32958134 -1.36365568  0.345306381  0.193055107  0.21721852
#> Macey       -0.19728082 -0.26927772 -0.363219506  0.368260269  0.21249474
#> Warners     -0.55673300 -0.26739400 -0.109470797  0.180283071  0.24208420
#> Zsivoczky   -0.11383190  0.03991159  0.538039776  0.585966156 -0.14271715
#> Hernu       -0.33168404  0.44308686  0.247293566  0.066908586 -0.20868256
#> Bernard     -0.32806564  0.36357920  0.006165316  0.279488675  0.32067773
#> Schwarzl    -0.68891640  0.56568604 -0.687053339 -0.008358849 -0.30211546
#> Pogorelov   -1.59379599  0.78370119 -0.037623661 -0.130531397 -0.03697576
#> Schoenbeck  0.16172381  0.85752368 -0.255850722  0.564222295  0.29680481
#> Barras      1.17415412  0.94512710  0.365550568  0.102255763  0.61186706

```

```
res.ind$contrib      # Contributions to the PCs
```

```

#>          Dim.1      Dim.2      Dim.3      Dim.4      Dim.5
#> SEBRLE      0.03854254  5.7118249  1.385418e+00  0.03572215  8.091161e+00
#> CLAY        0.65814114 13.8541889  6.460097e+00  8.55568792  4.034555e+00
#> BERNARD     1.86273218  6.1441319  1.349983e-01 19.57827284  4.202070e-02
#> YURKOV      0.79686310  0.4431309  2.147558e+01  2.57939100  9.859373e-01
#> ZSIVOCZKY   0.01178829  9.6816398  5.974848e+00  0.05231437  2.405750e-01
#> McMULLEN    0.01481737  2.3253860  2.496789e+00  9.13531719  1.558646e+01
#> MARTINEAU   6.03367104  3.9044125  2.830527e+00  0.73858431  3.118936e+01
#> HERNU       3.76996156  3.2842176  8.583863e-02  4.21505626  9.955149e-02
#> BARRAS      3.31942012  0.4024544  1.519980e+00  0.27758505  3.388731e-01
#> NOOL        8.12988880  5.8489726  1.293851e+00 12.83761115  1.257025e+01
#> BOURGUIGNON 20.53729577  3.7757623  3.534995e-04  2.35877858  1.088816e+00
#> Sebrle      12.55838616  3.5020697  9.881482e+00  3.47006223  3.889859e-03
#> Clay        11.59361384  6.2315181  1.341828e+00  0.53250375  2.007648e+00
#> Karpov      17.46609555  1.5072627  3.614914e+00  0.52101693  3.940874e+00
#> Macey       3.60207087 10.0732890  3.364879e+00  1.81387486  2.885677e-07
#> Warners     2.02910262  0.1879390  1.403071e+01  5.51585696  1.018550e-02
#> Zsivoczky   0.23441891  2.0310492  1.047894e+01  0.18126182  1.028128e+00

```

```

#> Hernu      0.08048777  3.2136178 1.020764e-01  3.82170515 2.946148e+00
#> Bernard    1.97075488  5.1661961 2.425213e+00  2.94737426 4.581507e+00
#> Schwarzl    0.53184785  1.0252129 3.800546e+00  4.56119277 5.167449e-01
#> Pogorelov   0.04843819  1.7533304 3.107757e-01  9.75034337 1.404313e-03
#> Schoenbeck  0.25864068  3.9969003 3.722687e-02  0.32970059 2.580092e+00
#> Barras      0.10519467  1.5876667 2.605305e+00  1.84296038 3.767994e+00
#>           Dim.6      Dim.7      Dim.8      Dim.9      Dim.10
#> SEBRLE      2.21256620  0.621426384 2.992045e-02 12.177477305 0.03819185
#> CLAY        17.48801877  0.651413899 6.035125e+00 0.101262442 3.58568943
#> BERNARD      1.65019840  2.373652810 7.847747e+00 0.810319793 0.34994507
#> YURKOV       0.10878629  1.516564073 2.094806e-01 0.336009790 0.51072064
#> ZSIVOCZKY    8.24561722  0.111923276 6.485544e-01 7.664919832 4.32741147
#> McMULLEN     0.35788945  3.287016354 1.360753e+00 0.312501167 5.51053518
#> MARTINEAU    2.14409841  1.291109482 1.347216e+00 1.398195851 13.54402896
#> HERNU        1.95655942  4.604850849 3.705288e+01 0.001071345 10.92781554
#> BARRAS       0.08376135  0.670038259 7.380544e+00 12.547331617 10.04537028
#> NOOL         2.54127369  4.067669683 2.975270e+00 0.375477289 0.33003418
#> BOURGUIGNON  0.06913582  5.144247534 5.032108e+00 7.744571086 0.12223626
#> Sebrle      0.07047579  0.001483775 1.481898e-02 20.105546253 1.72063803
#> Clay        6.20972751  0.952824148 5.340583e+00 4.703566841 13.52708188
#> Karpov       1.11680500 26.720158115 1.888802e+00 1.043988269 1.68193477
#> Macey        0.40014909  1.041910483 2.089853e+00 3.798767930 1.60957713
#> Warners      3.18673563  1.027384225 1.898339e-01 0.910422384 2.08904756
#> Zsivoczky    0.13322327  0.022889042 4.585705e+00 9.617852173 0.72605208
#> Hernu        1.13110069  2.821027418 9.687304e-01 0.125399768 1.55234328
#> Bernard      1.10655655  1.899449022 6.021268e-04 2.188071254 3.66566729
#> Schwarzl     4.87961053  4.598122119 7.477531e+00 0.001957159 3.25357879
#> Pogorelov    26.11665608  8.825322559 2.242329e-02 0.477268755 0.04873597
#> Schoenbeck   0.26890572 10.566272800 1.036933e+00 8.917302863 3.14020004
#> Barras       14.17432302 12.835417603 2.116763e+00 0.292892746 13.34533825

```

```
res.ind$cos2      # Quality of representation
```

```

#>           Dim.1      Dim.2      Dim.3      Dim.4      Dim.5
#> SEBRLE      0.007530179 0.49747323 8.132523e-02 0.001386688 2.689027e-01
#> CLAY        0.048701249 0.45701660 1.436281e-01 0.125791741 5.078506e-02
#> BERNARD      0.197199804 0.28996555 4.294015e-03 0.411819183 7.567259e-04
#> YURKOV       0.096109800 0.02382571 7.782303e-01 0.061812637 2.022798e-02
#> ZSIVOCZKY    0.001574385 0.57641944 2.397542e-01 0.001388216 5.465497e-03
#> McMULLEN     0.002175437 0.15219499 1.101379e-01 0.266486530 3.892621e-01
#> MARTINEAU    0.404013915 0.11654676 5.694575e-02 0.009826320 3.552552e-01
#> HERNU        0.399282749 0.15506199 2.731529e-03 0.088699901 1.793538e-03
#> BARRAS       0.616241975 0.03330700 8.478249e-02 0.010239088 1.070152e-02
#> NOOL         0.489872515 0.15711146 2.342405e-02 0.153694675 1.288433e-01
#> BOURGUIGNON  0.859698130 0.07045912 4.446015e-06 0.019618511 7.753120e-03
#> Sebrle      0.675380606 0.08395940 1.596674e-01 0.037079012 3.558507e-05
#> Clay        0.687592867 0.16475409 2.391051e-02 0.006274965 2.025440e-02
#> Karpov       0.783666922 0.03014772 4.873187e-02 0.004644764 3.007790e-02
#> Macey        0.363436037 0.45308203 1.020057e-01 0.036362957 4.952707e-09
#> Warners      0.255651956 0.01055582 5.311341e-01 0.138081100 2.182965e-04
#> Zsivoczky    0.045053176 0.17401353 6.051030e-01 0.006921739 3.361236e-02
#> Hernu        0.024824321 0.44184663 9.459148e-03 0.234196727 1.545686e-01
#> Bernard      0.289347476 0.33813318 1.069834e-01 0.085980212 1.144234e-01
#> Schwarzl     0.116721435 0.10030142 2.506043e-01 0.198892209 1.929118e-02
#> Pogorelov    0.007803472 0.12591966 1.504272e-02 0.312101619 3.848427e-05

```



```

#> Schoenbeck 0.067070098 0.46204603 2.900467e-03 0.016987442 1.138116e-01
#> Barras      0.018972684 0.12765099 1.411800e-01 0.066043061 1.156018e-01
#>
#>           Dim.6      Dim.7      Dim.8      Dim.9
#> SEBRLE      0.0443241299 8.907507e-03 3.890334e-04 8.954067e-02
#> CLAY        0.1326907339 3.536548e-03 2.972084e-02 2.820119e-04
#> BERNARD     0.0179131165 1.843634e-02 5.529104e-02 3.228572e-03
#> YURKOV      0.0013453555 1.341980e-02 1.681440e-03 1.525225e-03
#> ZSIVOCZKY   0.1129176906 1.096685e-03 5.764478e-03 3.852703e-02
#> McMULLEN    0.0053876990 3.540616e-02 1.329562e-02 1.726733e-03
#> MARTINEAU   0.0147210347 6.342774e-03 6.003515e-03 3.523552e-03
#> HERNU       0.0212478795 3.578167e-02 2.611676e-01 4.270425e-06
#> BARRAS      0.0015944528 9.126203e-03 9.118662e-02 8.766746e-02
#> NOOL        0.0157010551 1.798232e-02 1.193105e-02 8.514912e-04
#> BOURGUIGNON 0.0002967459 1.579887e-02 1.401866e-02 1.220108e-02
#> Sebrle      0.0003886276 5.854423e-06 5.303795e-05 4.069384e-02
#> Clay        0.0377627839 4.145976e-03 2.107924e-02 1.049876e-02
#> Karpov      0.0051379747 8.795817e-02 5.639959e-03 1.762907e-03
#> Macey       0.0041397727 7.712721e-03 1.403282e-02 1.442502e-02
#> Warners     0.0411689767 9.496848e-03 1.591742e-03 4.317040e-03
#> Zsivoczky   0.0026253777 3.227467e-04 5.865332e-02 6.956790e-02
#> Hernu       0.0357707217 6.383462e-02 1.988402e-02 1.455601e-03
#> Bernard     0.0166586433 2.046050e-02 5.883405e-06 1.209056e-02
#> Schwarzl    0.1098063093 7.403638e-02 1.092132e-01 1.616543e-05
#> Pogorelov   0.4314162233 1.043115e-01 2.404103e-04 2.893750e-03
#> Schoenbeck  0.0071500829 2.010275e-01 1.789520e-02 8.702893e-02
#> Barras      0.2621297474 1.698426e-01 2.540745e-02 1.988116e-03
#>
#>           Dim.10
#> SEBRLE      0.0002206741
#> CLAY        0.0078471026
#> BERNARD     0.0010956493
#> YURKOV      0.0018217256
#> ZSIVOCZKY   0.0170924251
#> McMULLEN    0.0239268142
#> MARTINEAU   0.0268211980
#> HERNU       0.0342288717
#> BARRAS      0.0551531863
#> NOOL        0.0005881295
#> BOURGUIGNON 0.0001513277
#> Sebrle      0.0027366539
#> Clay        0.0237264222
#> Karpov      0.0022318265
#> Macey       0.0048028954
#> Warners     0.0077841113
#> Zsivoczky   0.0041268259
#> Hernu       0.0141595965
#> Bernard     0.0159167991
#> Schwarzl    0.0211173850
#> Pogorelov   0.0002322016
#> Schoenbeck  0.0240826922
#> Barras      0.0711836486

```


Chapter 3

Predict using PCA

In this section, we'll show how to predict the coordinates of supplementary individuals and variables using only the information provided by the previously performed PCA.

1. Data: rows 24 to 27 and columns 1 to 10 [in decathlon2 data sets]. The new data must contain columns (variables) with the same names and in the same order as the active data used to compute PCA.

```
# Data for the supplementary individuals
ind.sup <- decathlon2[24:27, 1:10]
ind.sup[, 1:6]
```

```
#>      X100m Long.jump Shot.put High.jump X400m X110m.hurdle
#> KARPOV  11.02      7.30   14.77      2.04 48.37      14.09
#> WARNERS 11.11      7.60   14.31      1.98 48.68      14.23
#> Nool    10.80      7.53   14.26      1.88 48.81      14.80
#> Drews   10.87      7.38   13.07      1.88 48.51      14.01
```

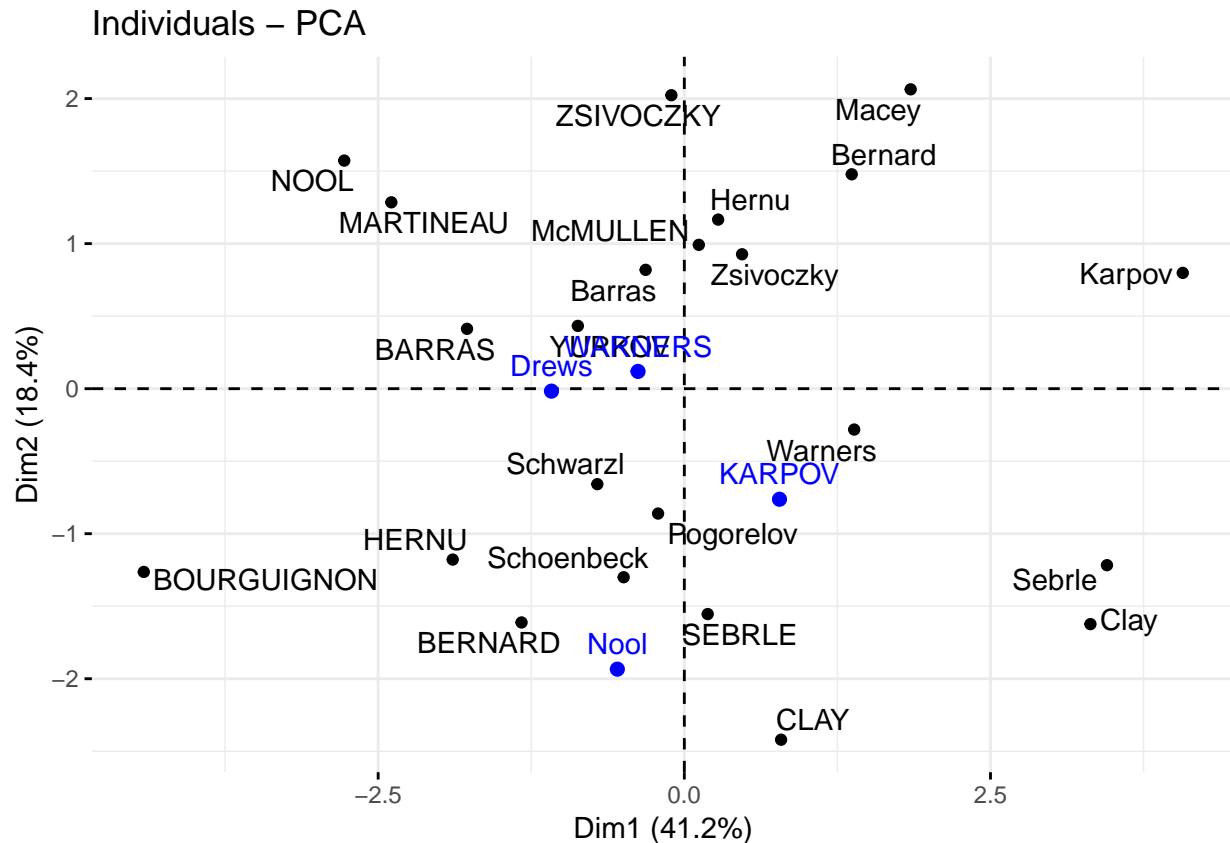
2. Predict the coordinates of new individuals data. Use the R base function predict():

```
ind.sup.coord <- predict(res.pca, newdata = ind.sup)
ind.sup.coord[, 1:4]
```

```
#>      PC1      PC2      PC3      PC4
#> KARPOV  0.7772521 -0.76237804 1.5971253 1.6863286
#> WARNERS -0.3779697 0.11891968 1.7005146 -0.6908084
#> Nool    -0.5468405 -1.93402211 0.4724184 -2.2283706
#> Drews   -1.0848227 -0.01703198 2.9818031 -1.5006207
```

3. Graph of individuals including the supplementary individuals:

```
# Plot of active individuals
p <- fviz_pca_ind(res.pca, repel = TRUE)
# Add supplementary individuals
fviz_add(p, ind.sup.coord, color = "blue")
```



The predicted coordinates of individuals can be manually calculated as follow:

1. Center and scale the new individuals data using the center and the scale of the PCA
2. Calculate the predicted coordinates by multiplying the scaled values with the eigenvectors (loadings) of the principal components. The R code below can be used :

```
# Centering and scaling the supplementary individuals
ind.scaled <- scale(ind.sup,
                    center = res.pca$center,
                    scale = res.pca$scale)

# Coordinates of the individuals
coord_func <- function(ind, loadings){
  r <- loadings*ind
  apply(r, 2, sum)
}

pca.loadings <- res.pca$rotation
ind.sup.coord <- t(apply(ind.scaled, 1, coord_func, pca.loadings ))
ind.sup.coord[, 1:4]
```

```
#>           PC1      PC2      PC3      PC4
#> KARPPOV  0.7772521 -0.76237804 1.5971253  1.6863286
#> WARNERS -0.3779697  0.11891968 1.7005146 -0.6908084
#> Nool     -0.5468405 -1.93402211 0.4724184 -2.2283706
#> Drews    -1.0848227 -0.01703198 2.9818031 -1.5006207
```

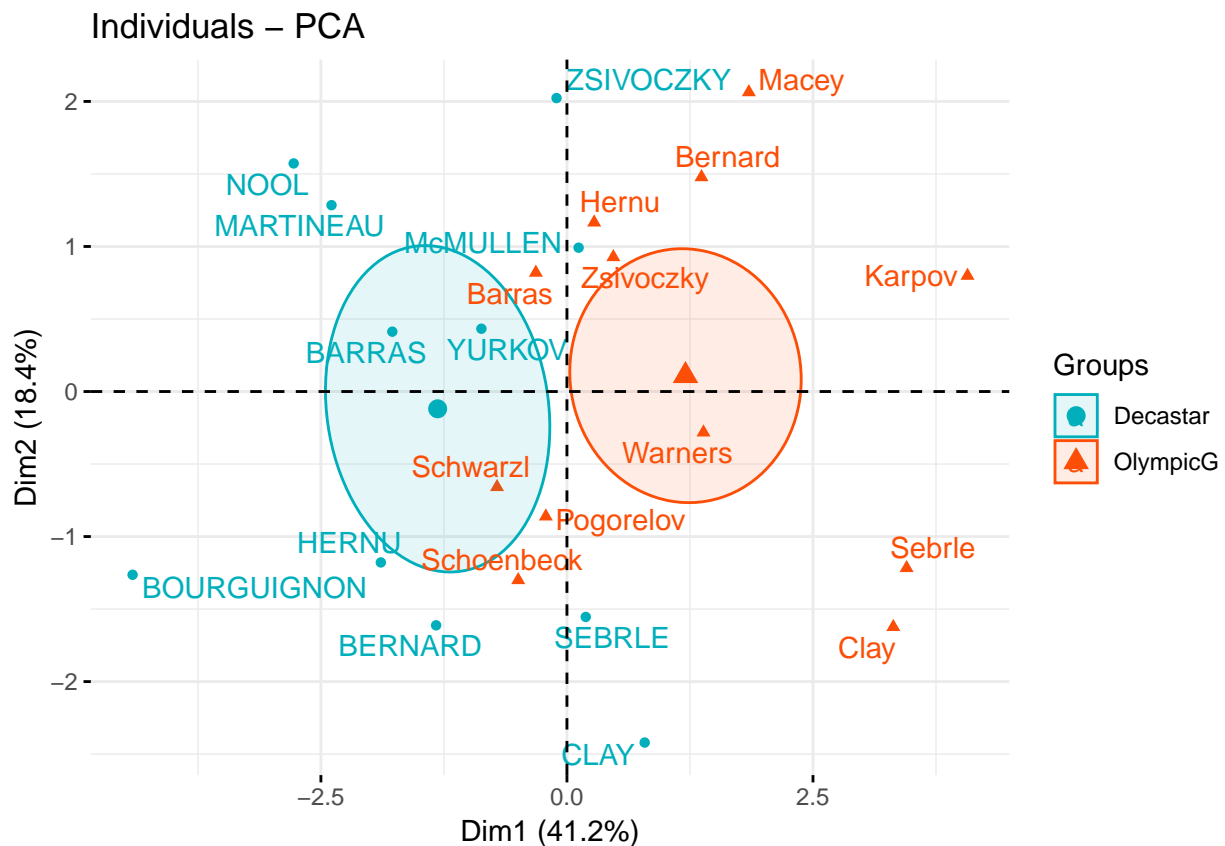
3.1 Supplementary variables

3.1.1 Qualitative / categorical variables

The data sets decathlon2 contain a supplementary qualitative variable at columns 13 corresponding to the type of competitions.

Qualitative / categorical variables can be used to color individuals by groups. The grouping variable should be of same length as the number of active individuals (here 23).

```
groups <- as.factor(decathlon2$Competition[1:23])
fviz_pca_ind(res.pca,
  col.ind = groups, # color by groups
  palette = c("#00AFBB", "#FC4E07"),
  addEllipses = TRUE, # Concentration ellipses
  ellipse.type = "confidence",
  legend.title = "Groups",
  repel = TRUE
)
```



Calculate the coordinates for the levels of grouping variables. The coordinates for a given group is calculated as the mean coordinates of the individuals in the group.

```
library(magrittr) # for pipe %>%
library(dplyr)   # everything else

# 1. Individual coordinates
res.ind <- get_pca_ind(res.pca)
```

```
# 2. Coordinate of groups
coord.groups <- res.ind$coord %>%
  as_data_frame() %>%
  select(Dim.1, Dim.2) %>%
  mutate(competition = groups) %>%
  group_by(competition) %>%
  summarise(
    Dim.1 = mean(Dim.1),
    Dim.2 = mean(Dim.2)
  )
coord.groups
```

```
#> # A tibble: 2 x 3
#>   competition Dim.1 Dim.2
#>   <fct>         <dbl> <dbl>
#> 1 Decastar     -1.31 -0.119
#> 2 OlympicG      1.20  0.109
```

3.1.2 Quantitative variables

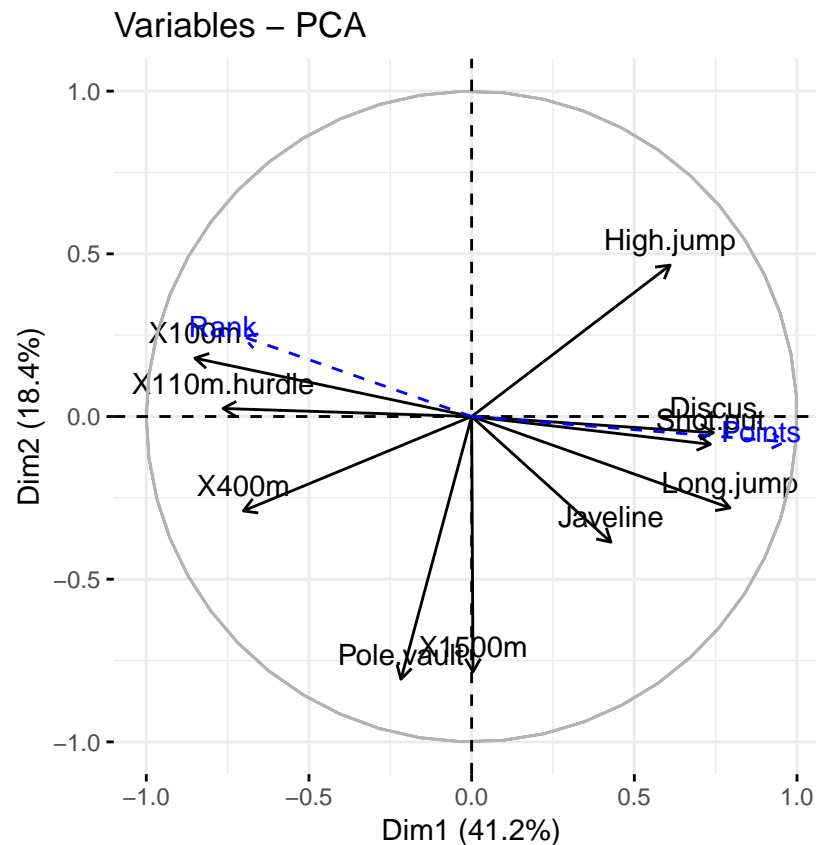
Data: columns 11:12. Should be of same length as the number of active individuals (here 23)

```
quanti.sup <- decathlon2[1:23, 11:12, drop = FALSE]
head(quanti.sup)
```

```
#>           Rank Points
#> SEBRLE      1    8217
#> CLAY        2    8122
#> BERNARD     4    8067
#> YURKOV      5    8036
#> ZSIVOCZKY   7    8004
#> McMULLEN    8    7995
```

The coordinates of a given quantitative variable are calculated as the correlation between the quantitative variables and the principal components.

```
# Predict coordinates and compute cos2
quanti.coord <- cor(quanti.sup, res.pca$x)
quanti.cos2 <- quanti.coord^2
# Graph of variables including supplementary variables
p <- fviz_pca_var(res.pca)
fviz_add(p, quanti.coord, color = "blue", geom = "arrow")
```



3.2 Theory behind PCA results

3.2.1 PCA results for variables

Here we'll show how to calculate the PCA results for variables: coordinates, cos2 and contributions:

`var.coord` = loadings * the component standard deviations `var.cos2` = `var.coord`² `var.contrib`. The contribution of a variable to a given principal component is (in percentage) : (`var.cos2` * 100) / (total cos2 of the component)

```
# Helper function
#::::::::::::::::::::::::::::::::::::
var_coord_func <- function(loadings, comp.sdev){
  loadings*comp.sdev
}

# Compute Coordinates
#::::::::::::::::::::::::::::::::::::
loadings <- res.pca$rotation
sdev <- res.pca$sdev
var.coord <- t(apply(loadings, 1, var_coord_func, sdev))
head(var.coord[, 1:4])
```

```
#>          PC1      PC2      PC3      PC4
#> X100m    -0.8506257  0.17939806 -0.3015564  0.03357320
#> Long.jump  0.7941806 -0.28085695  0.1905465 -0.11538956
```

```
#> Shot.put      0.7339127 -0.08540412 -0.5175978  0.12846837
#> High.jump     0.6100840  0.46521415 -0.3300852  0.14455012
#> X400m         -0.7016034 -0.29017826 -0.2835329  0.43082552
#> X110m.hurdle -0.7641252  0.02474081 -0.4488873 -0.01689589
```

```
# Compute Cos2
```

```
#::::::::::::::::::::::::::::::::::::
```

```
var.cos2 <- var.coord^2
head(var.cos2[, 1:4])
```

```
#>          PC1          PC2          PC3          PC4
#> X100m      0.7235641 0.0321836641 0.09093628 0.0011271597
#> Long.jump  0.6307229 0.0788806285 0.03630798 0.0133147506
#> Shot.put   0.5386279 0.0072938636 0.26790749 0.0165041211
#> High.jump  0.3722025 0.2164242070 0.10895622 0.0208947375
#> X400m      0.4922473 0.0842034209 0.08039091 0.1856106269
#> X110m.hurdle 0.5838873 0.0006121077 0.20149984 0.0002854712
```

```
# Compute contributions
```

```
#::::::::::::::::::::::::::::::::::::
```

```
comp.cos2 <- apply(var.cos2, 2, sum)
contrib <- function(var.cos2, comp.cos2){var.cos2*100/comp.cos2}
var.contrib <- t(apply(var.cos2,1, contrib, comp.cos2))
head(var.contrib[, 1:4])
```

```
#>          PC1          PC2          PC3          PC4
#> X100m      17.544293  1.7505098  7.338659  0.13755240
#> Long.jump  15.293168  4.2904162  2.930094  1.62485936
#> Shot.put   13.060137  0.3967224 21.620432  2.01407269
#> High.jump   9.024811 11.7715838  8.792888  2.54987951
#> X400m      11.935544  4.5799296  6.487636 22.65090599
#> X110m.hurdle 14.157544  0.0332933 16.261261  0.03483735
```

3.2.2 PCA results for individuals

- `ind.coord = res.pca$x`
- Cos2 of individuals. Two steps:
 - Calculate the square distance between each individual and the PCA center of gravity: $d2 = [(var1_ind_i - mean_var1)/sd_var1]^2 + \dots + [(var10_ind_i - mean_var10)/sd_var10]^2 + \dots + \dots$
 - Calculate the cos2 as $ind.coord^2/d2$
- Contributions of individuals to the principal components: $100 * (1 / number_of_individuals) * (ind.coord^2 / comp_sdev^2)$. Note that the sum of all the contributions per column is 100

```
# Coordinates of individuals
```

```
#::::::::::::::::::::::::::::::::::::
```

```
ind.coord <- res.pca$x
head(ind.coord[, 1:4])
```

```
#>          PC1          PC2          PC3          PC4
#> SEBRLE     0.1912074 -1.5541282 -0.6283688  0.08205241
#> CLAY       0.7901217 -2.4204156  1.3568870  1.26984296
#> BERNARD    -1.3292592 -1.6118687 -0.1961500 -1.92092203
#> YURKOV     -0.8694134  0.4328779 -2.4739822  0.69723814
#> ZSIVOCZKY -0.1057450  2.0233632  1.3049312 -0.09929630
#> McMULLEN   0.1185550  0.9916237  0.8435582  1.31215266
```



```

# Cos2 of individuals
#::::::::::::::::::::::::::::::::::
# 1. square of the distance between an individual and the
# PCA center of gravity
center <- res.pca$center
scale<- res.pca$scale

getdistance <- function(ind_row, center, scale){
  return(sum(((ind_row-center)/scale)^2))
}

d2 <- apply(decathlon2.active,1, getdistance, center, scale)
# 2. Compute the cos2. The sum of each row is 1
cos2 <- function(ind.coord, d2){return(ind.coord^2/d2)}
ind.cos2 <- apply(ind.coord, 2, cos2, d2)
head(ind.cos2[, 1:4])

```

```

#>           PC1      PC2      PC3      PC4
#> SEBRLE  0.007530179 0.49747323 0.081325232 0.001386688
#> CLAY    0.048701249 0.45701660 0.143628117 0.125791741
#> BERNARD 0.197199804 0.28996555 0.004294015 0.411819183
#> YURKOV  0.096109800 0.02382571 0.778230322 0.061812637
#> ZSIVOCZKY 0.001574385 0.57641944 0.239754152 0.001388216
#> McMULLEN 0.002175437 0.15219499 0.110137872 0.266486530

```

```

# Contributions of individuals
#::::::::::::::::::::::::::::::::::
contrib <- function(ind.coord, comp.sdev, n.ind){
  100*(1/n.ind)*ind.coord^2/comp.sdev^2
}
ind.contrib <- t(apply(ind.coord, 1, contrib,
  res.pca$sdev, nrow(ind.coord)))
head(ind.contrib[, 1:4])

```

```

#>           PC1      PC2      PC3      PC4
#> SEBRLE  0.03854254  5.7118249  1.3854184  0.03572215
#> CLAY    0.65814114 13.8541889  6.4600973  8.55568792
#> BERNARD 1.86273218  6.1441319  0.1349983 19.57827284
#> YURKOV  0.79686310  0.4431309 21.4755770  2.57939100
#> ZSIVOCZKY 0.01178829 9.6816398  5.9748485  0.05231437
#> McMULLEN 0.01481737 2.3253860  2.4967890  9.13531719

```


Chapter 4

Diagnostic Plots

I have tried to use `fortify` function in `ggplot2` which can access different statistics related to linear model. The basic diagnostic plot which we often get using `plot` function in the fitted model using `lm` command.

The function `diagPlots` gives an list of six different plots which can be arranged in a grid using `grid` and `gridExtra` packages.

```
library(ggplot2)

diagPlot<-function(model){
  p1<-ggplot(model, aes(.fitted, .resid))+geom_point()
  p1<-p1+stat_smooth(method="loess")+geom_hline(yintercept=0, col="red", linetype="dashed")
  p1<-p1+xlab("Fitted values")+ylab("Residuals")
  p1<-p1+ggtitle("Residual vs Fitted Plot")+theme_bw()

  p2<-ggplot(model, aes(qqnorm(.stdresid)[[1]], .stdresid))+geom_point(na.rm = TRUE)
  p2<-p2+geom_abline(aes(qqline(.stdresid)))+xlab("Theoretical Quantiles")+ylab("Standardized Residuals")
  p2<-p2+ggtitle("Normal Q-Q")+theme_bw()

  p3<-ggplot(model, aes(.fitted, sqrt(abs(.stdresid))))+geom_point(na.rm=TRUE)
  p3<-p3+stat_smooth(method="loess", na.rm = TRUE)+xlab("Fitted Value")
  p3<-p3+ylab(expression(sqrt("|Standardized residuals|")))
  p3<-p3+ggtitle("Scale-Location")+theme_bw()

  p4<-ggplot(model, aes(seq_along(.cooks), .cooks))+geom_bar(stat="identity", position="identity")
  p4<-p4+xlab("Obs. Number")+ylab("Cook's distance")
  p4<-p4+ggtitle("Cook's distance")+theme_bw()

  p5<-ggplot(model, aes(.hat, .stdresid))+geom_point(aes(size=.cooks), na.rm=TRUE)
  p5<-p5+stat_smooth(method="loess", na.rm=TRUE)
  p5<-p5+xlab("Leverage")+ylab("Standardized Residuals")
  p5<-p5+ggtitle("Residual vs Leverage Plot")
  p5<-p5+scale_size_continuous("Cook's Distance", range=c(1,5))
  p5<-p5+theme_bw()+theme(legend.position="bottom")

  p6<-ggplot(model, aes(.hat, .cooks))+geom_point(na.rm=TRUE)+stat_smooth(method="loess", na.rm=TRUE)
  p6<-p6+xlab("Leverage hii")+ylab("Cook's Distance")
  p6<-p6+ggtitle("Cook's dist vs Leverage hii/(1-hii)")
  p6<-p6+geom_abline(slope=seq(0,3,0.5), color="gray", linetype="dashed")
}
```

```

p6<-p6+theme_bw()

return(list(rvfPlot=p1, qqPlot=p2, sclLocPlot=p3, cdPlot=p4, rvlevPlot=p5, cvlPlot=p6))
}

```

Using the mtcars datasets, a linear model is fitted with mpg as response and cyl, disp, hp, drat and wt has predictor variable

```

lm.model <- lm(mpg ~ cyl+disp+hp+drat+wt, data=mtcars)
diagPlts <- diagPlot(lm.model)

```

To display the plots in a grid, some packages mentioned above should be installed.

```

lbry <- c("grid", "gridExtra")
lapply(lbry, require, character.only=TRUE, warn.conflicts = FALSE, quietly = TRUE)

```

```

:> [[1]]
:> [1] TRUE
:>
:> [[2]]
:> [1] TRUE

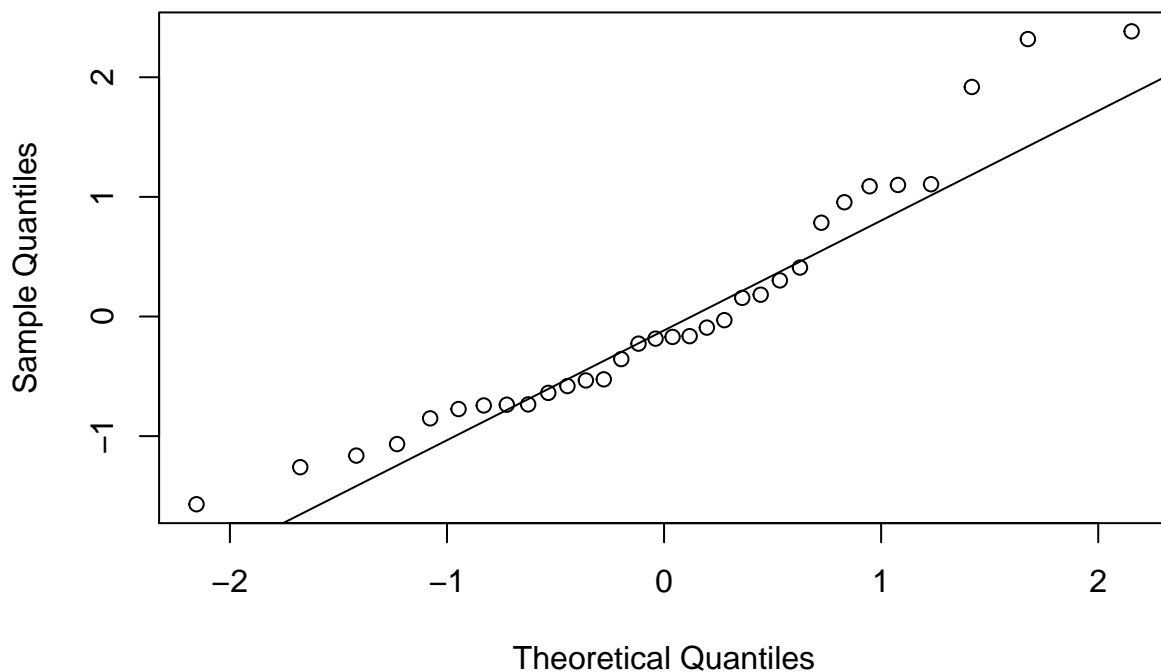
```

Thus the plot obtained is,

```
do.call(grid.arrange, c(diagPlts, main="Diagnostic Plots", ncol=3))
```

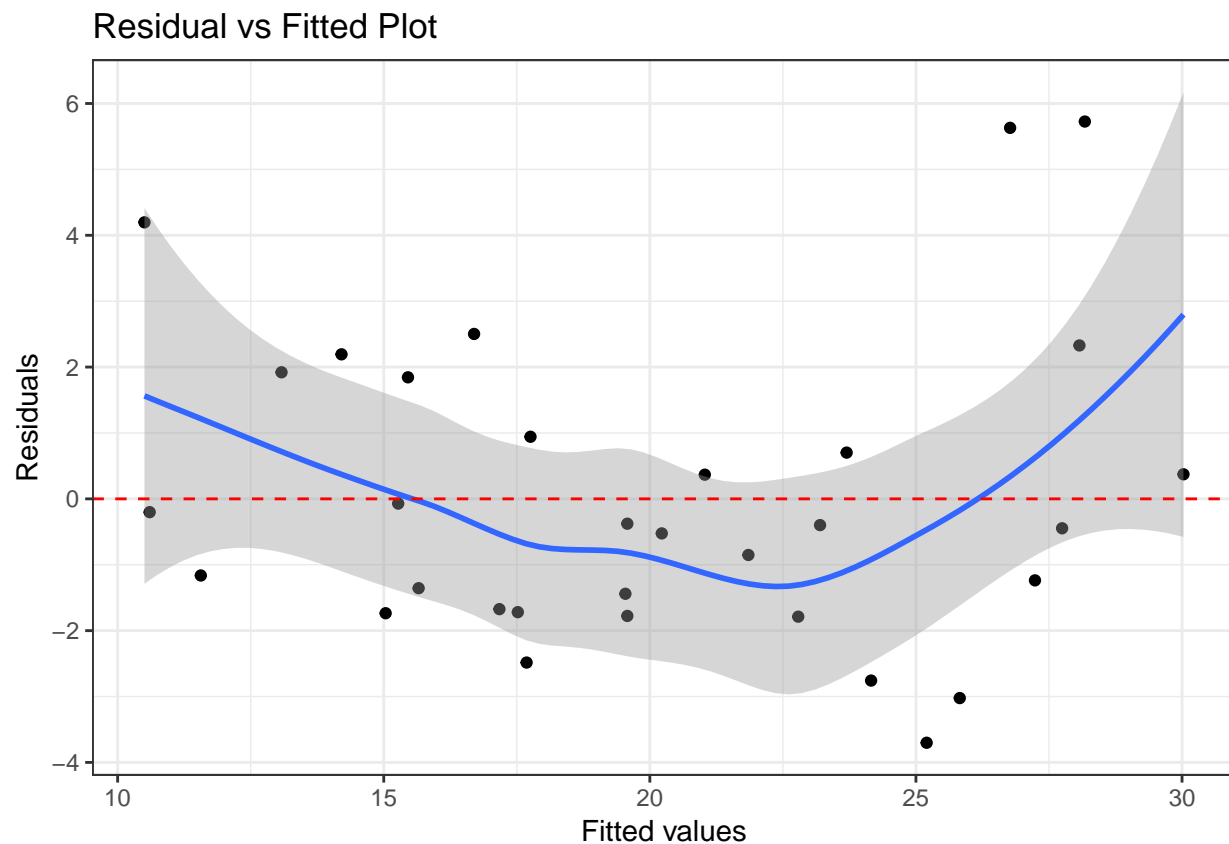
```
:> Error: Aesthetics must be either length 1 or the same as the data (32): x
```

Normal Q–Q Plot



4.1 Residual vs Fitted plot

```
model <- lm.model
p1<-ggplot(model, aes(.fitted, .resid)) +
  geom_point() +
  stat_smooth(method="loess") +
  geom_hline(yintercept=0, col="red", linetype="dashed") +
  xlab("Fitted values") +
  ylab("Residuals") +
  ggtitle("Residual vs Fitted Plot") +
  theme_bw()
p1
```

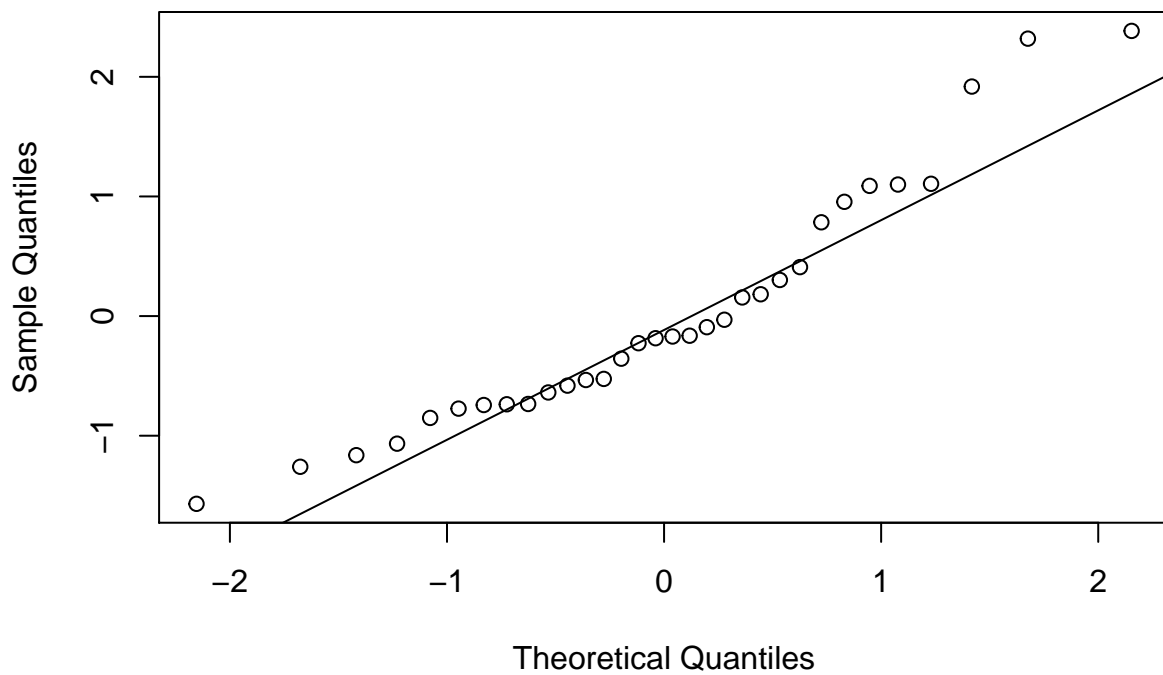


4.2 Normal QQ

```
p2 <- ggplot(model, aes(qqnorm(.stdresid)[[1]], .stdresid)) +
  geom_point(na.rm = TRUE) +
  geom_abline(aes(qqline(.stdresid)))
# xlab("Theoretical Quantiles") +
# ylab("Standardized Residuals") +
# ggtitle("Normal Q-Q") +
# theme_bw()
p2
```

```
:> Error: Aesthetics must be either length 1 or the same as the data (32): x
```

Normal Q-Q Plot

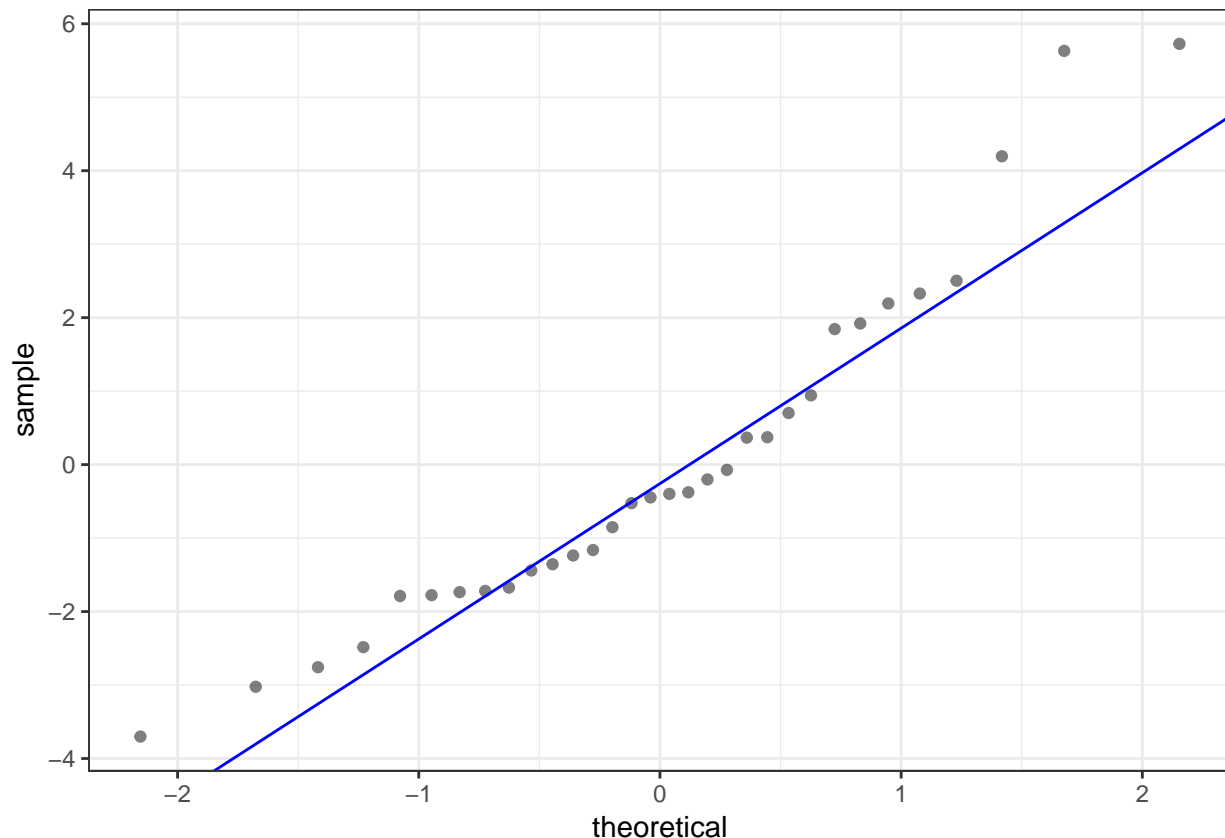


```

y <- quantile(model$resid[!is.na(model$resid)], c(0.25, 0.75))
x <- qnorm(c(0.25, 0.75))
slope <- diff(y)/diff(x)
int <- y[1L] - slope * x[1L]
p <- ggplot(model, aes(sample = .resid)) +
  stat_qq(alpha = 0.5) +
  geom_abline(slope = slope, intercept = int, color="blue") +
  theme_bw()

```

p



The standard Q-Q diagnostic for linear models plots quantiles of the standardized residuals vs. theoretical quantiles of $N(0,1)$. Peter's `ggQQ` function plots the residuals. The snippet below amends that and adds a few cosmetic changes to make the plot more like what one would get from `plot(lm(...))`.

```

# https://stackoverflow.com/a/19990107/5270873
ggQQ = function(lm) {
  # extract standardized residuals from the fit
  d <- data.frame(std.resid = rstandard(lm))
  # calculate 1Q/4Q line
  y <- quantile(d$std.resid[!is.na(d$std.resid)], c(0.25, 0.75))
  x <- qnorm(c(0.25, 0.75))
  slope <- diff(y)/diff(x)
  int <- y[1L] - slope * x[1L]

  p <- ggplot(data=d, aes(sample=std.resid)) +
    stat_qq(shape=1, size=3) + # open circles

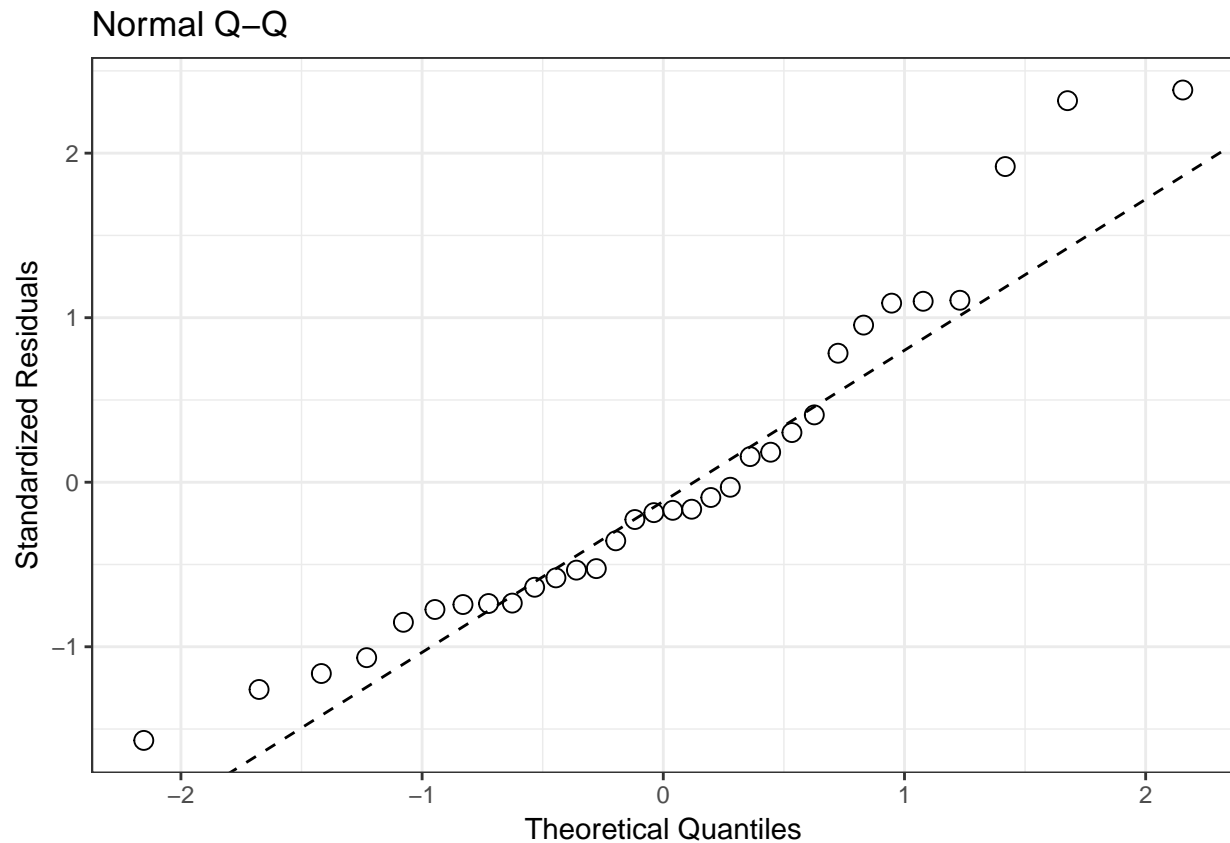
```

```

labs(title="Normal Q-Q",           # plot title
      x="Theoretical Quantiles",   # x-axis label
      y="Standardized Residuals") + # y-axis label
geom_abline(slope = slope, intercept = int, linetype="dashed") + # dashed reference line
theme_bw()
return(p)
}

ggQQ(model)

```



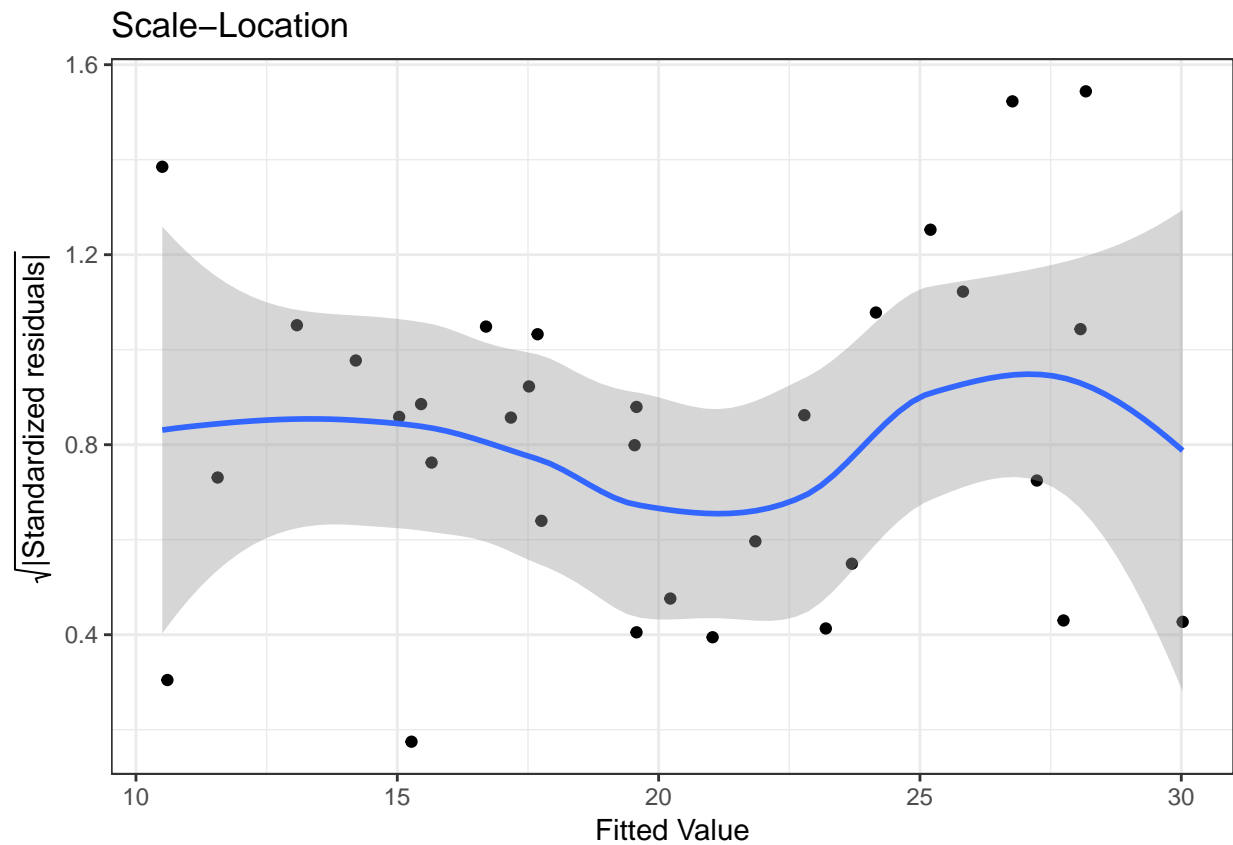
4.3 Scale-location

```

p3 <- ggplot(model, aes(.fitted, sqrt(abs(.stdresid)))) +
  geom_point(na.rm=TRUE) +
  stat_smooth(method="loess", na.rm = TRUE) +
  xlab("Fitted Value") +
  ylab(expression(sqrt("|Standardized residuals|"))) +
  ggtitle("Scale-Location") +
  theme_bw()

p3

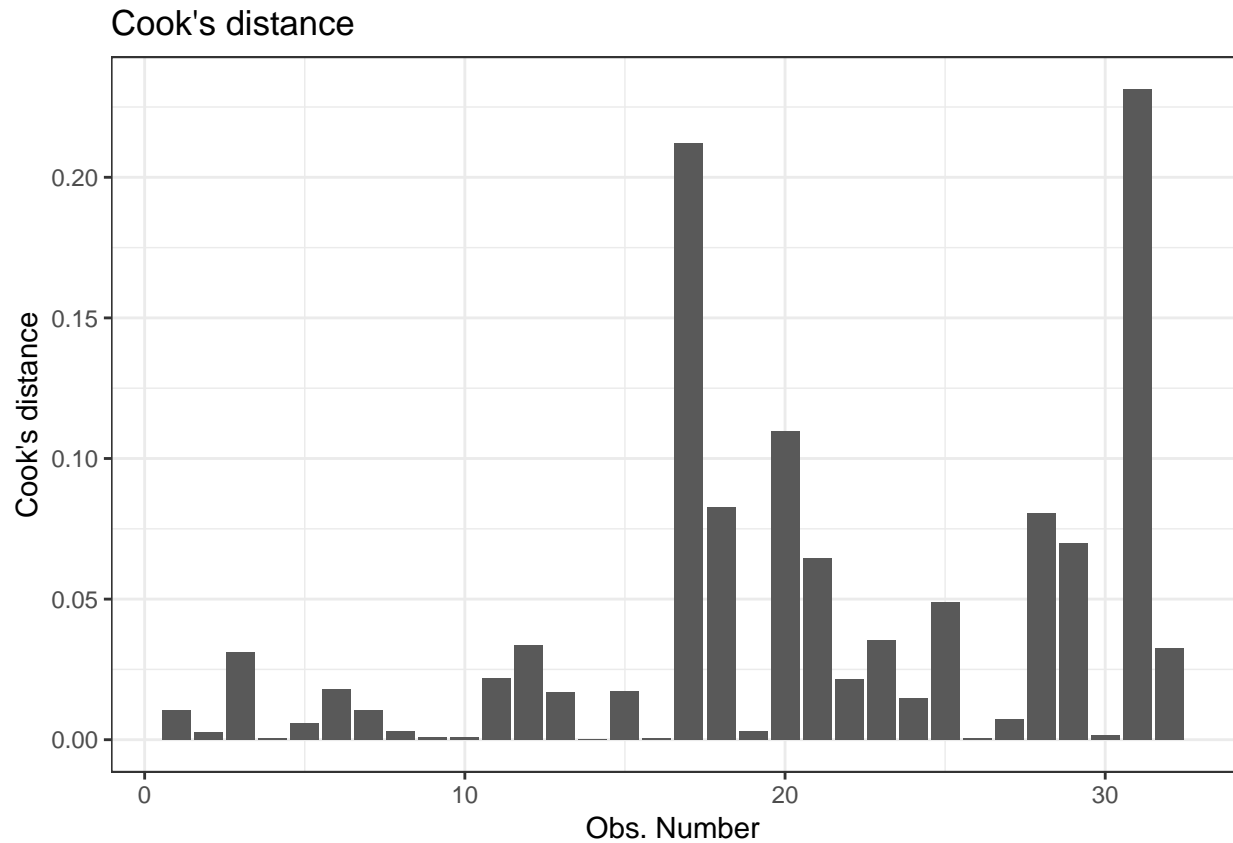
```

4.4 Cook's Distance

```
p4 <- ggplot(model, aes(seq_along(.cooks_d), .cooks_d)) +
  geom_bar(stat="identity", position="identity") +
  xlab("Obs. Number") +
  ylab("Cook's distance") +
  ggtitle("Cook's distance") +
  theme_bw()
```

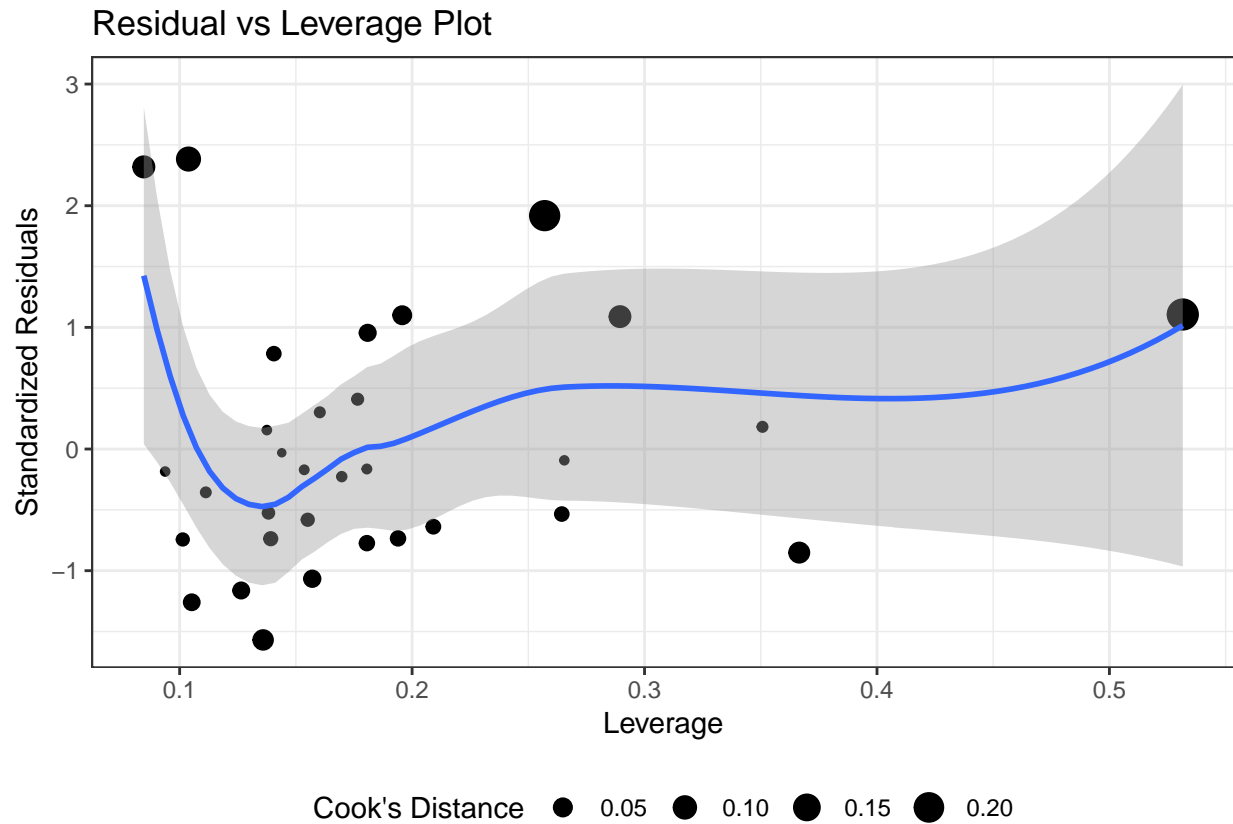
p4



4.5 Residual vs Leverage Plot

```
p5 <- ggplot(model, aes(.hat, .stdresid)) +
  geom_point(aes(size=.cooks), na.rm=TRUE) +
  stat_smooth(method="loess", na.rm=TRUE) +
  xlab("Leverage") +
  ylab("Standardized Residuals") +
  ggtitle("Residual vs Leverage Plot") +
  scale_size_continuous("Cook's Distance", range=c(1,5)) +
  theme_bw() +
  theme(legend.position="bottom")
```

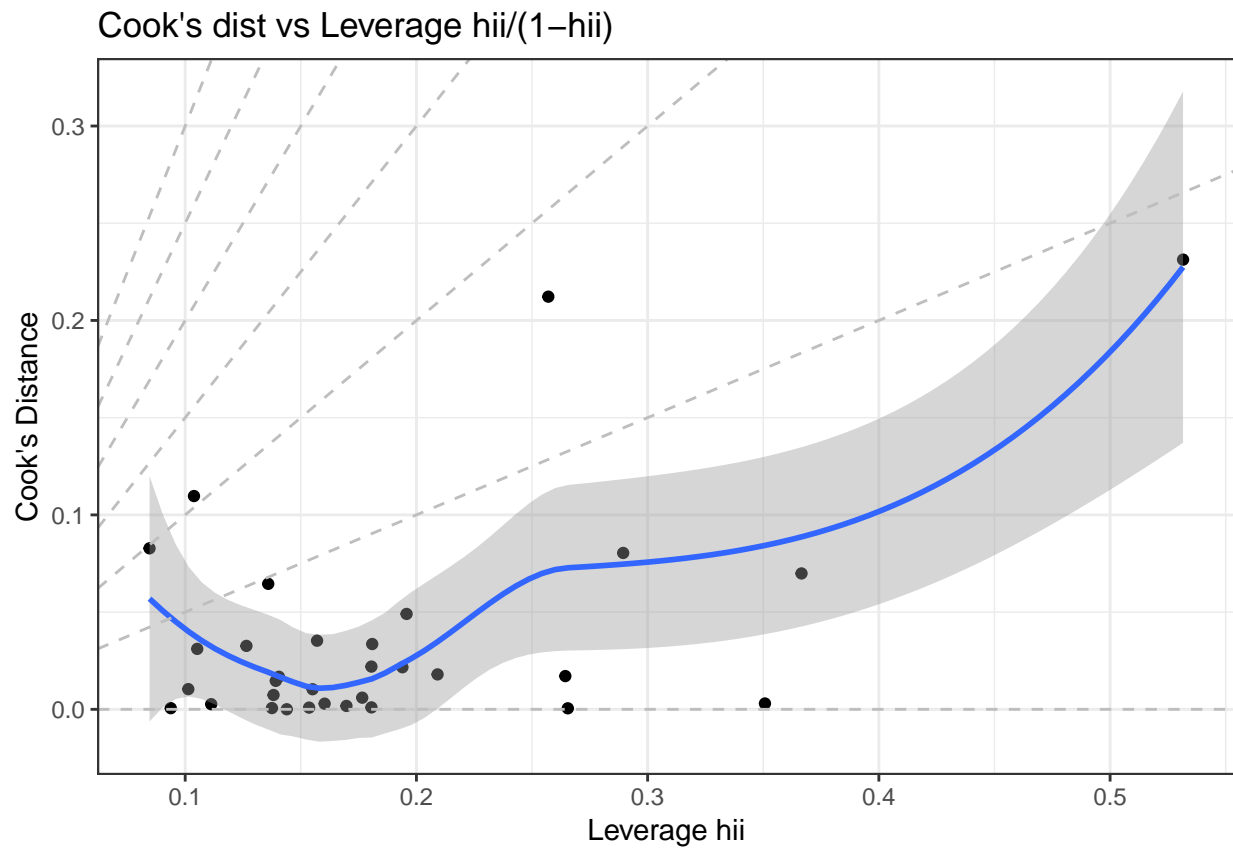
p5



4.6 Cook's dist vs Leverage $h_{ii}/(1-h_{ii})$

```
p6 <- ggplot(model, aes(.hat, .cooksdist)) +
  geom_point(na.rm=TRUE) +
  stat_smooth(method="loess", na.rm=TRUE) +
  xlab("Leverage hii") +
  ylab("Cook's Distance") +
  ggtitle("Cook's dist vs Leverage hii/(1-hii)") +
  geom_abline(slope=seq(0,3,0.5), color="gray", linetype="dashed") +
  theme_bw()
```

p6



Chapter 5

Temperature modeling using nested dataframes

5.1 Prepare the data

[http://ijlyttle.github.io/isugg_purrr/presentation.html#\(1\)](http://ijlyttle.github.io/isugg_purrr/presentation.html#(1))

5.1.1 Packages to run this presentation

```
library("readr")
library("tibble")
library("dplyr")
library("tidyr")
library("stringr")
library("ggplot2")
library("purrr")
library("broom")
```

5.1.2 Motivation

As you know, purrr is a recent package from Hadley Wickham, focused on lists and functional programming, like dplyr is focused on data-frames.

I figure a good way to learn a new package is to try to solve a problem, so we have a dataset:

- you can view or download
- you can download the source of this presentation
- these are three temperatures recorded simultaneously in a piece of electronics
- it will be very valuable to be able to characterize the transient temperature for each sensor
- we want to apply the same set of models across all three sensors
- it will be easier to show using pictures

5.1.3 Let's get the data into shape

Using the readr package

```
temperature_wide <-
  read_csv(file.path(data_raw_dir, "temperature.csv")) %>%
  print()

# A tibble: 327 x 4
   instant      temperature_a temperature_b temperature_c
  <dtm>          <dbl>          <dbl>          <dbl>
1 2015-11-13 06:10:19      116.          91.7          84.2
2 2015-11-13 06:10:23      116.          91.7          84.2
3 2015-11-13 06:10:27      116.          91.6          84.2
4 2015-11-13 06:10:31      116.          91.7          84.2
5 2015-11-13 06:10:36      116.          91.7          84.2
6 2015-11-13 06:10:41      116.          91.6          84.2
7 2015-11-13 06:10:46      116.          91.5          84.2
8 2015-11-13 06:10:51      116.          91.5          84.2
9 2015-11-13 06:10:56      116.          91.5          84.2
10 2015-11-13 06:11:01      115.          91.5          84.2
# ... with 317 more rows
```

5.1.4 Is temperature_wide “tidy”?

```
# A tibble: 327 x 4
   instant      temperature_a temperature_b temperature_c
  <dtm>          <dbl>          <dbl>          <dbl>
1 2015-11-13 06:10:19      116.          91.7          84.2
2 2015-11-13 06:10:23      116.          91.7          84.2
3 2015-11-13 06:10:27      116.          91.6          84.2
4 2015-11-13 06:10:31      116.          91.7          84.2
5 2015-11-13 06:10:36      116.          91.7          84.2
6 2015-11-13 06:10:41      116.          91.6          84.2
7 2015-11-13 06:10:46      116.          91.5          84.2
8 2015-11-13 06:10:51      116.          91.5          84.2
9 2015-11-13 06:10:56      116.          91.5          84.2
10 2015-11-13 06:11:01      115.          91.5          84.2
# ... with 317 more rows
```

Why or why not?

5.1.5 Tidy data

1. Each column is a variable
2. Each row is an observation
3. Each cell is a value

(<http://www.jstatsoft.org/v59/i10/paper>)

My personal observation is that “tidy” can depend on the context, on what you want to do with the data.

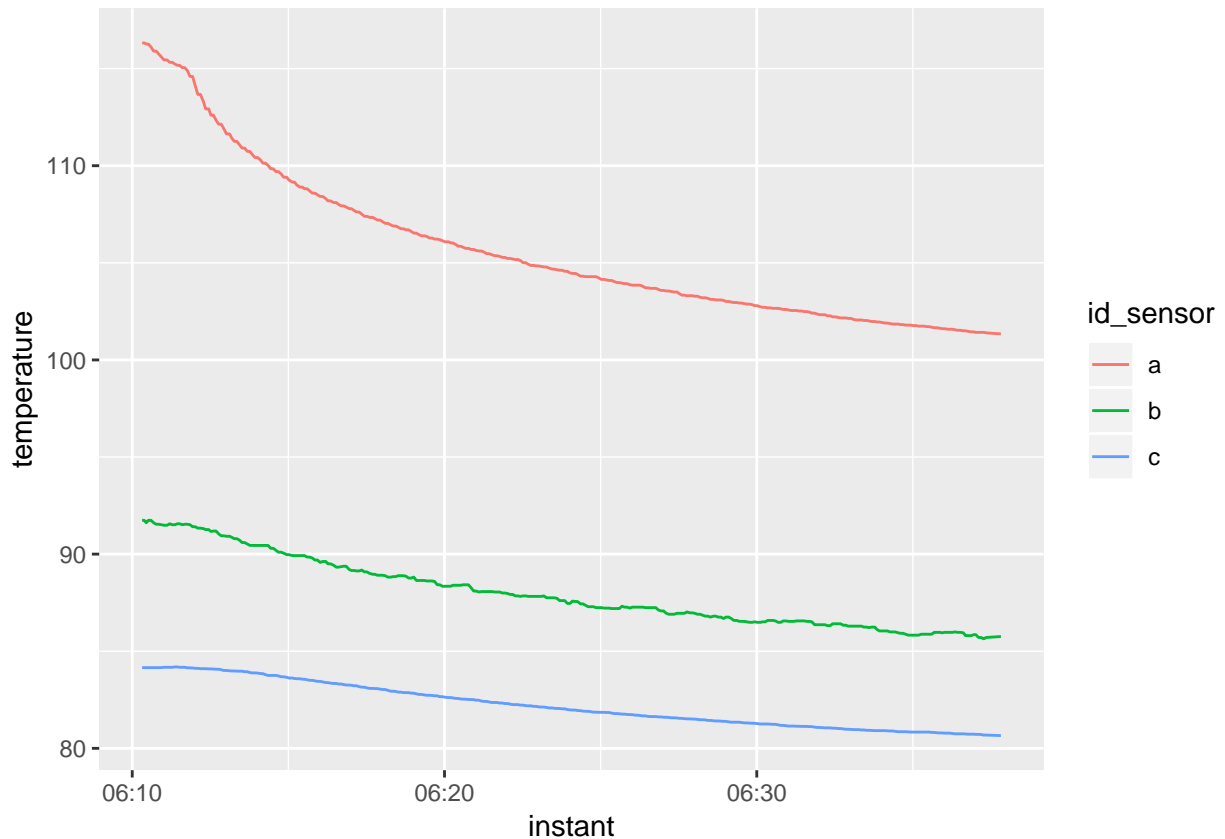
5.1.6 Let's get this into a tidy form

```
temperature_tall <-
  temperature_wide %>%
  gather(key = "id_sensor", value = "temperature", starts_with("temp")) %>%
  mutate(id_sensor = str_replace(id_sensor, "temperature_", "")) %>%
  print()
```

```
# A tibble: 981 x 3
  instant          id_sensor temperature
  <dtm>           <chr>         <dbl>
1 2015-11-13 06:10:19 a             116.
2 2015-11-13 06:10:23 a             116.
3 2015-11-13 06:10:27 a             116.
4 2015-11-13 06:10:31 a             116.
5 2015-11-13 06:10:36 a             116.
6 2015-11-13 06:10:41 a             116.
7 2015-11-13 06:10:46 a             116.
8 2015-11-13 06:10:51 a             116.
9 2015-11-13 06:10:56 a             116.
10 2015-11-13 06:11:01 a             115.
# ... with 971 more rows
```

5.1.7 Now, it's easier to visualize

```
temperature_tall %>%
  ggplot(aes(x = instant, y = temperature, color = id_sensor)) +
  geom_line()
```



5.1.8 Calculate delta time (Δt) and delta temperature (ΔT)

delta_time Δt

change in time since event started, s

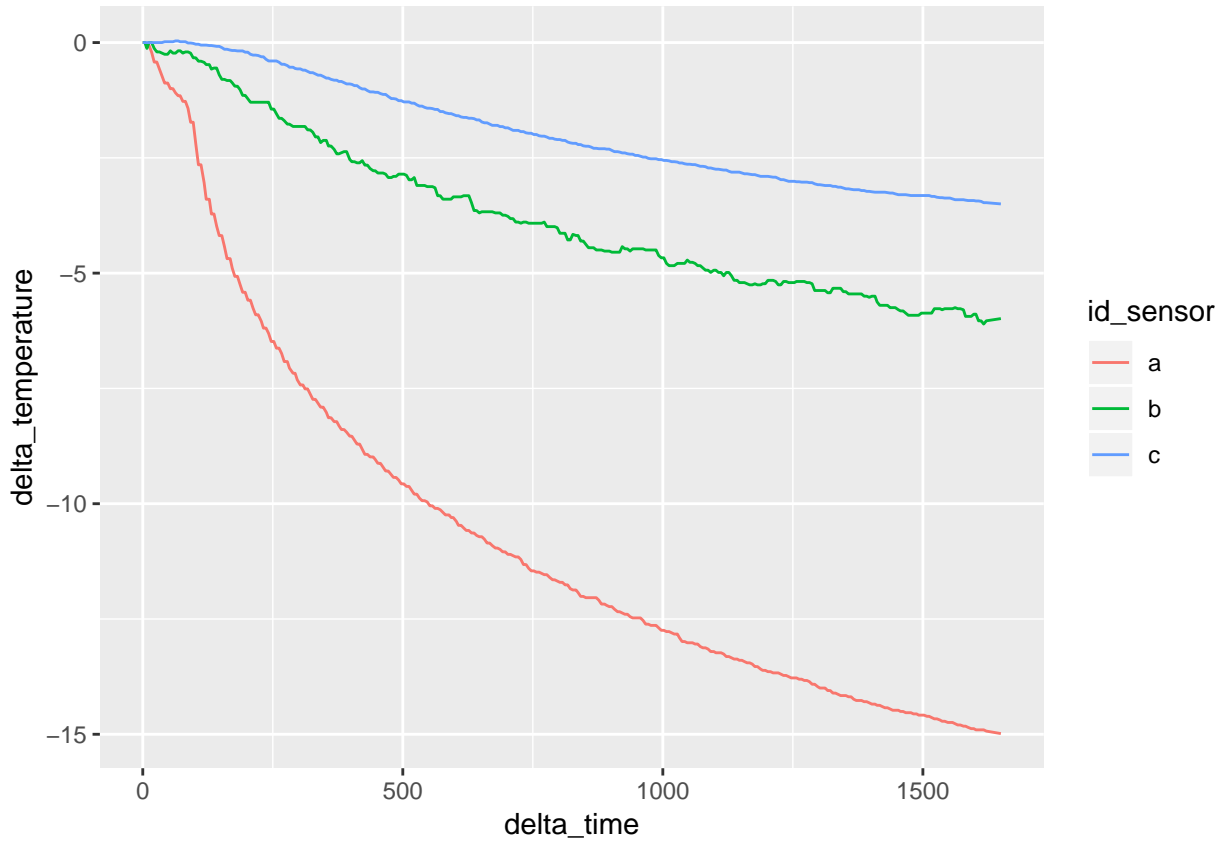
delta_temperature: ΔT

change in temperature since event started, °C

```
delta <-
  temperature_tall %>%
  arrange(id_sensor, instant) %>%
  group_by(id_sensor) %>%
  mutate(
    delta_time = as.numeric(instant) - as.numeric(instant[[1]]),
    delta_temperature = temperature - temperature[[1]]
  ) %>%
  select(id_sensor, delta_time, delta_temperature)
```

5.1.9 Let's have a look

```
# plot delta time vs delta temperature, by sensor
delta %>%
  ggplot(aes(x = delta_time, y = delta_temperature, color = id_sensor)) +
  geom_line()
```

5.2 Define the models

We want to see how three different curve-fits might perform on these three data-sets:

5.2.0.1 Newtonian cooling

$$\Delta T = \Delta T_0 * (1 - e^{-\frac{\delta t}{\tau_0}})$$

5.2.1 Semi-infinite solid

$$\Delta T = \Delta T_0 * \operatorname{erfc}\left(\sqrt{\frac{\tau_0}{\delta t}}\right)$$

5.2.2 Semi-infinite solid with convection

$$\Delta T = \Delta T_0 * \left[\operatorname{erfc}\left(\sqrt{\frac{\tau_0}{\delta t}}\right) - e^{Bi_0 + \left(\frac{Bi_0}{2}\right)^2 \frac{\delta t}{\tau_0}} * \operatorname{erfc}\left(\sqrt{\frac{\tau_0}{\delta t}} + \frac{Bi_0}{2} * \sqrt{\frac{\delta t}{\tau_0}}\right) \right]$$

5.2.3 erf and erfc functions

```
# reference: http://stackoverflow.com/questions/29067916/r-error-function-erfz
# (see Abramowitz and Stegun 29.2.29)
erf <- function(x) 2 * pnorm(x * sqrt(2)) - 1
erfc <- function(x) 2 * pnorm(x * sqrt(2), lower = FALSE)
```

5.2.4 Newton cooling equation

```
newton_cooling <- function(x) {
  nls(
    delta_temperature ~ delta_temperature_0 * (1 - exp(-delta_time/tau_0)),
    start = list(delta_temperature_0 = -10, tau_0 = 50),
    data = x
  )
}
```

5.2.5 Temperature models: simple and convection

```
semi_infinite_simple <- function(x) {
  nls(
    delta_temperature ~ delta_temperature_0 * erfc(sqrt(tau_0 / delta_time)),
    start = list(delta_temperature_0 = -10, tau_0 = 50),
    data = x
  )
}

semi_infinite_convection <- function(x){
  nls(
    delta_temperature ~
      delta_temperature_0 * (
        erfc(sqrt(tau_0 / delta_time)) -
        exp(Bi_0 + (Bi_0/2)^2 * delta_time / tau_0) *
        erfc(sqrt(tau_0 / delta_time) +
          (Bi_0/2) * sqrt(delta_time / tau_0))
      ),
    start = list(delta_temperature_0 = -5, tau_0 = 50, Bi_0 = 1.e6),
    data = x
  )
}
```

5.3 Test modeling on one dataset

5.3.1 Before going into purrr

Before doing anything, we want to show that we can do something with one dataset and one model-function:

```
# only one sensor; it is a test
tmp_data <- delta %>% filter(id_sensor == "a")

tmp_model <- newton_cooling(tmp_data)
```

```
summary(tmp_model)
```

```
Formula: delta_temperature ~ delta_temperature_0 * (1 - exp(-delta_time/tau_0))
```

```
Parameters:
```

```

              Estimate Std. Error t value Pr(>|t|)
delta_temperature_0 -15.06085    0.05262  -286.2   <2e-16 ***
tau_0               500.01382    4.83673   103.4   <2e-16 ***
---

```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.3267 on 325 degrees of freedom
```

```
Number of iterations to convergence: 7
```

```
Achieved convergence tolerance: 4.136e-06
```

5.3.2 Look at predictions

```

# apply prediction and make it tidy
tmp_pred <-
  tmp_data %>%
  mutate(modeled = predict(tmp_model, data = .)) %>%
  select(id_sensor, delta_time, measured = delta_temperature, modeled) %>%
  gather("type", "delta_temperature", measured:modeled) %>%
  print()

```

```

# A tibble: 654 x 4
# Groups:   id_sensor [1]
  id_sensor delta_time type      delta_temperature
  <chr>      <dbl> <chr>      <dbl>
1 a          0 measured          0
2 a          4 measured          0
3 a          8 measured        -0.06
4 a         12 measured        -0.06
5 a         17 measured       -0.211
6 a         22 measured       -0.423
7 a         27 measured       -0.423
8 a         32 measured       -0.574
9 a         37 measured       -0.726
10 a        42 measured       -0.878
# ... with 644 more rows

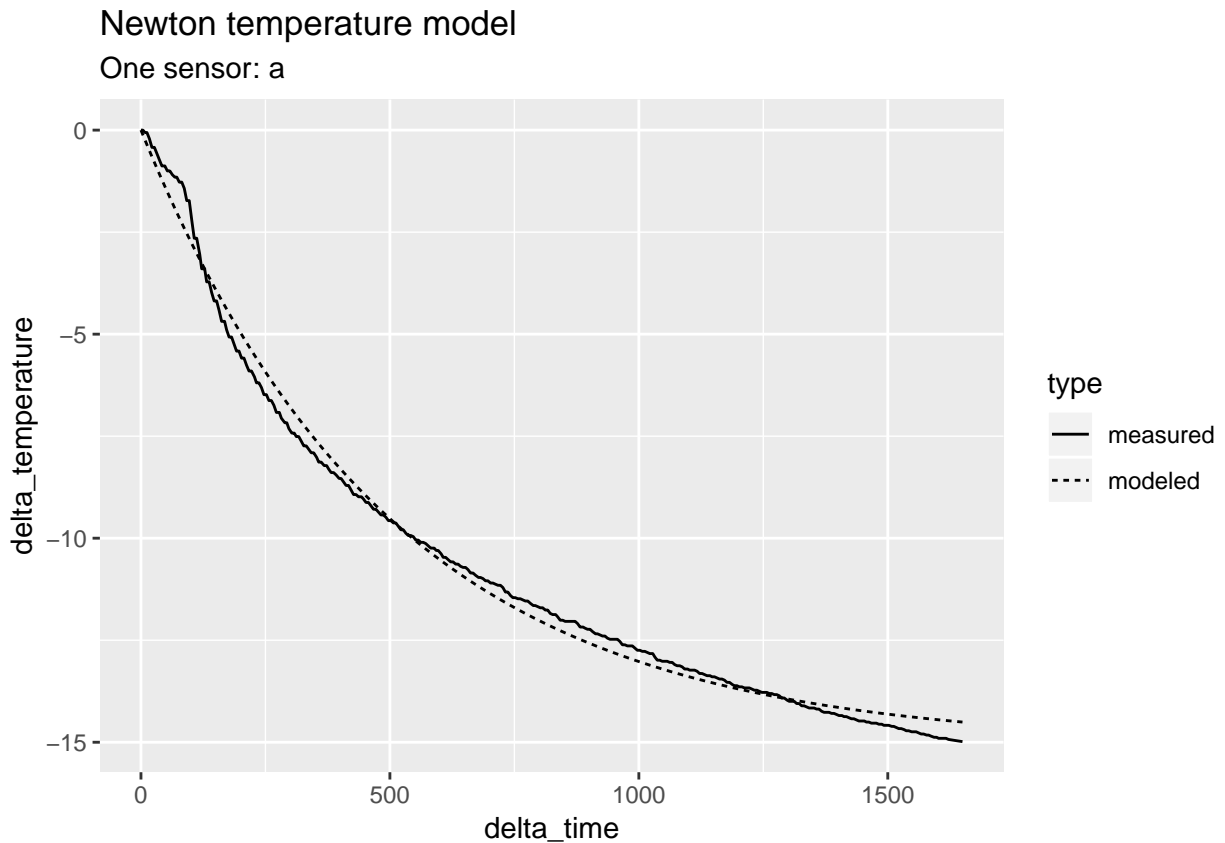
```

5.3.3 Plot Newton model

```

tmp_pred %>%
  ggplot(aes(x = delta_time, y = delta_temperature, linetype = type)) +
  geom_line() +
  labs(title = "Newton temperature model", subtitle = "One sensor: a")

```



5.3.4 “Regular” data-frame (deltas)

```
print(delta)
```

```
# A tibble: 981 x 3
# Groups:   id_sensor [3]
  id_sensor delta_time delta_temperature
  <chr>      <dbl>      <dbl>
1 a          0          0
2 a          4          0
3 a          8        -0.06
4 a         12        -0.06
5 a         17       -0.211
6 a         22       -0.423
7 a         27       -0.423
8 a         32       -0.574
9 a         37       -0.726
10 a        42       -0.878
# ... with 971 more rows
```

Each column of the dataframe is a vector - in this case, a character vector and two doubles

5.4 Making a nested dataframe

5.4.1 How to make a weird data-frame

Here's where the fun starts - a column of a data-frame can be a list.

- use `tidyr::nest()` to makes a column `data`, which is a list of data-frames
- this seems like a stronger expression of the `dplyr::group_by()` idea

```
# nest delta_time and delta_temperature variables
```

```
delta_nested <-
  delta %>%
  nest(-id_sensor) %>%
  print()
```

```
# A tibble: 3 x 2
  id_sensor data
  <chr>      <list>
1 a        <tibble [327 x 2]>
2 b        <tibble [327 x 2]>
3 c        <tibble [327 x 2]>
```

5.4.2 Map dataframes to a modeling function (Newton)

- `map()` is like `lapply()`
- `map()` returns a list-column (it keeps the weirdness)

```
model_nested <-
  delta_nested %>%
  mutate(model = map(data, newton_cooling)) %>%
  print()
```

```
# A tibble: 3 x 3
  id_sensor data          model
  <chr>      <list>      <list>
1 a        <tibble [327 x 2]> <nls>
2 b        <tibble [327 x 2]> <nls>
3 c        <tibble [327 x 2]> <nls>
```

We get an additional list-column `model`.

5.4.3 We can use `map2()` to make the predictions

- `map2()` is like `mapply()`
- designed to map two columns (`model`, `data`) to a function `predict()`

```
predict_nested <-
  model_nested %>%
  mutate(pred = map2(model, data, predict)) %>%
  print()
```

```
# A tibble: 3 x 4
  id_sensor data          model pred
  <chr>      <list>      <list> <list>
```

```

1 a      <tibble [327 x 2]> <nls>  <dbl [327]>
2 b      <tibble [327 x 2]> <nls>  <dbl [327]>
3 c      <tibble [327 x 2]> <nls>  <dbl [327]>

```

Another list-column `pred` for the prediction results.

5.4.4 We need to get out of the weirdness

- use `unnest()` to get back to a regular data-frame

```

predict_unnested <-
  predict_nested %>%
  unnest(data, pred) %>%
  print()

```

```

# A tibble: 981 x 4
  id_sensor  pred delta_time delta_temperature
  <chr>      <dbl>    <dbl>          <dbl>
1 a          0         0             0
2 a      -0.120         4             0
3 a      -0.239         8          -0.06
4 a      -0.357        12          -0.06
5 a      -0.503        17          -0.211
6 a      -0.648        22          -0.423
7 a      -0.792        27          -0.423
8 a      -0.934        32          -0.574
9 a      -1.07         37          -0.726
10 a     -1.21         42          -0.878
# ... with 971 more rows

```

5.4.5 We can wrangle the predictions

- get into a form that makes it easier to plot

```

predict_tall <-
  predict_unnested %>%
  rename(modeled = pred, measured = delta_temperature) %>%
  gather("type", "delta_temperature", modeled, measured) %>%
  print()

```

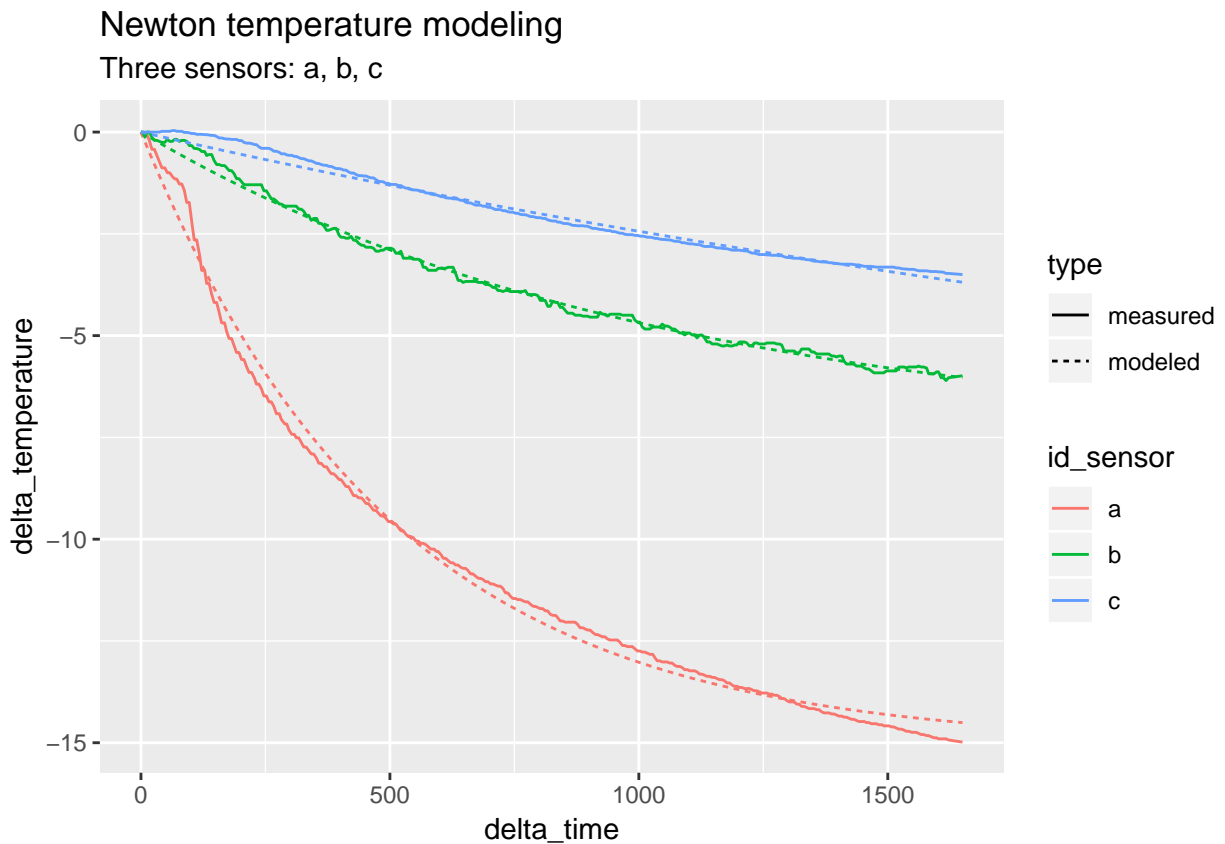
```

# A tibble: 1,962 x 4
  id_sensor delta_time type    delta_temperature
  <chr>      <dbl> <chr>          <dbl>
1 a          0 modeled          0
2 a          4 modeled      -0.120
3 a          8 modeled      -0.239
4 a         12 modeled      -0.357
5 a         17 modeled      -0.503
6 a         22 modeled      -0.648
7 a         27 modeled      -0.792
8 a         32 modeled      -0.934
9 a         37 modeled      -1.07
10 a        42 modeled      -1.21
# ... with 1,952 more rows

```

5.4.6 We can visualize the predictions

```
predict_tall %>%
  ggplot(aes(x = delta_time, y = delta_temperature)) +
  geom_line(aes(color = id_sensor, linetype = type)) +
  labs(title = "Newton temperature modeling",
       subtitle = "Three sensors: a, b, c")
```



5.5 Apply multiple models on a nested structure

5.5.1 Step 1: Selection of models

Make a list of functions to model:

```
list_model <-
  list(
    newton_cooling = newton_cooling,
    semi_infinite_simple = semi_infinite_simple,
    semi_infinite_convection = semi_infinite_convection
  )
```

5.5.2 Step 2: write a function to define the “inner” loop

```
# add additional variable with the model name

fn_model <- function(.model, df) {
  # one parameter for the model in the list, the second for the data
  # safer to avoid non-standard evaluation
  # df %>% mutate(model = map(data, .model))

  df$model <- map(df$data, possibly(.model, NULL))
  df
}
```

- for a given model-function and a given (weird) data-frame, return a modified version of that data-frame with a column `model`, which is the model-function applied to each element of the data-frame’s `data` column (which is itself a list of data-frames)
- the purrr functions `safely()` and `possibly()` are **very** interesting. I think they could be useful outside of purrr as a friendlier way to do error-handling.

5.5.3 Step 3: Use `map_df()` to define the “outer” loop

```
# this dataframe will be the second input of fn_model
delta_nested %>%
  print()

# A tibble: 3 x 2
  id_sensor data
  <chr>      <list>
1 a         <tibble [327 x 2]>
2 b         <tibble [327 x 2]>
3 c         <tibble [327 x 2]>

# fn_model is receiving two inputs: one from list_model and from delta_nested
model_nested_new <-
  list_model %>%
  map_df(fn_model, delta_nested, .id = "id_model") %>%
  print()
```

```
# A tibble: 9 x 4
  id_model      id_sensor data      model
  <chr>      <chr>      <list>      <list>
1 newton_cooling a         <tibble [327 x 2]> <nls>
2 newton_cooling b         <tibble [327 x 2]> <nls>
3 newton_cooling c         <tibble [327 x 2]> <nls>
4 semi_infinite_simple a         <tibble [327 x 2]> <nls>
5 semi_infinite_simple b         <tibble [327 x 2]> <nls>
6 semi_infinite_simple c         <tibble [327 x 2]> <nls>
7 semi_infinite_convection a         <tibble [327 x 2]> <NULL>
8 semi_infinite_convection b         <tibble [327 x 2]> <NULL>
9 semi_infinite_convection c         <tibble [327 x 2]> <NULL>
```

- for each element of a list of model-functions, run the inner-loop function, and row-bind the results into a data-frame

- we want to discard the rows where the model failed
- we also want to investigate why they failed, but that's a different talk

5.5.4 Step 4: Use `map()` to identify the null models

```
model_nested_new <-
  list_model %>%
  map_df(fn_model, delta_nested, .id = "id_model") %>%
  mutate(is_null = map(model, is.null)) %>%
  print()
```

```
# A tibble: 9 x 5
  id_model      id_sensor data          model is_null
  <chr>        <chr>    <list>      <list> <list>
1 newton_cooling a      <tibble [327 x 2]> <nls> <lgl [1]>
2 newton_cooling b      <tibble [327 x 2]> <nls> <lgl [1]>
3 newton_cooling c      <tibble [327 x 2]> <nls> <lgl [1]>
4 semi_infinite_simple a    <tibble [327 x 2]> <nls> <lgl [1]>
5 semi_infinite_simple b    <tibble [327 x 2]> <nls> <lgl [1]>
6 semi_infinite_simple c    <tibble [327 x 2]> <nls> <lgl [1]>
7 semi_infinite_convection a  <tibble [327 x 2]> <NULL> <lgl [1]>
8 semi_infinite_convection b  <tibble [327 x 2]> <NULL> <lgl [1]>
9 semi_infinite_convection c  <tibble [327 x 2]> <NULL> <lgl [1]>
```

- using `map(model, is.null)` returns a list column
- to use `filter()`, we have to escape the weirdness

5.5.5 Step 5: `map_lgl()` to identify nulls and get out of the weirdness

```
model_nested_new <-
  list_model %>%
  map_df(fn_model, delta_nested, .id = "id_model") %>%
  mutate(is_null = map_lgl(model, is.null)) %>%
  print()
```

```
# A tibble: 9 x 5
  id_model      id_sensor data          model is_null
  <chr>        <chr>    <list>      <list> <lgl>
1 newton_cooling a      <tibble [327 x 2]> <nls> FALSE
2 newton_cooling b      <tibble [327 x 2]> <nls> FALSE
3 newton_cooling c      <tibble [327 x 2]> <nls> FALSE
4 semi_infinite_simple a    <tibble [327 x 2]> <nls> FALSE
5 semi_infinite_simple b    <tibble [327 x 2]> <nls> FALSE
6 semi_infinite_simple c    <tibble [327 x 2]> <nls> FALSE
7 semi_infinite_convection a  <tibble [327 x 2]> <NULL> TRUE
8 semi_infinite_convection b  <tibble [327 x 2]> <NULL> TRUE
9 semi_infinite_convection c  <tibble [327 x 2]> <NULL> TRUE
```

- using `map_lgl(model, is.null)` returns a vector column

5.5.6 Step 6: filter() nulls and select() variables to clean up

```
model_nested_new <-
  list_model %>%
  map_df(fn_model, delta_nested, .id = "id_model") %>%
  mutate(is_null = map_lgl(model, is.null)) %>%
  filter(!is_null) %>%
  select(-is_null) %>%
  print()

# A tibble: 6 x 4
  id_model      id_sensor data      model
  <chr>        <chr>    <list>    <list>
1 newton_cooling a      <tibble [327 x 2]> <nls>
2 newton_cooling b      <tibble [327 x 2]> <nls>
3 newton_cooling c      <tibble [327 x 2]> <nls>
4 semi_infinite_simple a    <tibble [327 x 2]> <nls>
5 semi_infinite_simple b    <tibble [327 x 2]> <nls>
6 semi_infinite_simple c    <tibble [327 x 2]> <nls>
```

5.5.7 Step 7: Calculate predictions on nested dataframe

```
predict_nested <-
  model_nested_new %>%
  mutate(pred = map2(model, data, predict)) %>%
  print()

# A tibble: 6 x 5
  id_model      id_sensor data      model pred
  <chr>        <chr>    <list>    <list> <list>
1 newton_cooling a      <tibble [327 x 2]> <nls> <dbl [327]>
2 newton_cooling b      <tibble [327 x 2]> <nls> <dbl [327]>
3 newton_cooling c      <tibble [327 x 2]> <nls> <dbl [327]>
4 semi_infinite_simple a    <tibble [327 x 2]> <nls> <dbl [327]>
5 semi_infinite_simple b    <tibble [327 x 2]> <nls> <dbl [327]>
6 semi_infinite_simple c    <tibble [327 x 2]> <nls> <dbl [327]>
```

5.5.8 unnest(), make it tall and tidy

```
predict_tall <-
  predict_nested %>%
  unnest(data, pred) %>%
  rename(modeled = pred, measured = delta_temperature) %>%
  gather("type", "delta_temperature", modeled, measured) %>%
  print()

# A tibble: 3,924 x 5
  id_model      id_sensor delta_time type      delta_temperature
  <chr>        <chr>        <dbl> <chr>        <dbl>
1 newton_cooling a              0 modeled          0
2 newton_cooling a              4 modeled     -0.120
3 newton_cooling a              8 modeled     -0.239
```

```

4 newton_cooling a          12 modeled          -0.357
5 newton_cooling a          17 modeled          -0.503
6 newton_cooling a          22 modeled          -0.648
7 newton_cooling a          27 modeled          -0.792
8 newton_cooling a          32 modeled          -0.934
9 newton_cooling a          37 modeled          -1.07
10 newton_cooling a         42 modeled          -1.21
# ... with 3,914 more rows

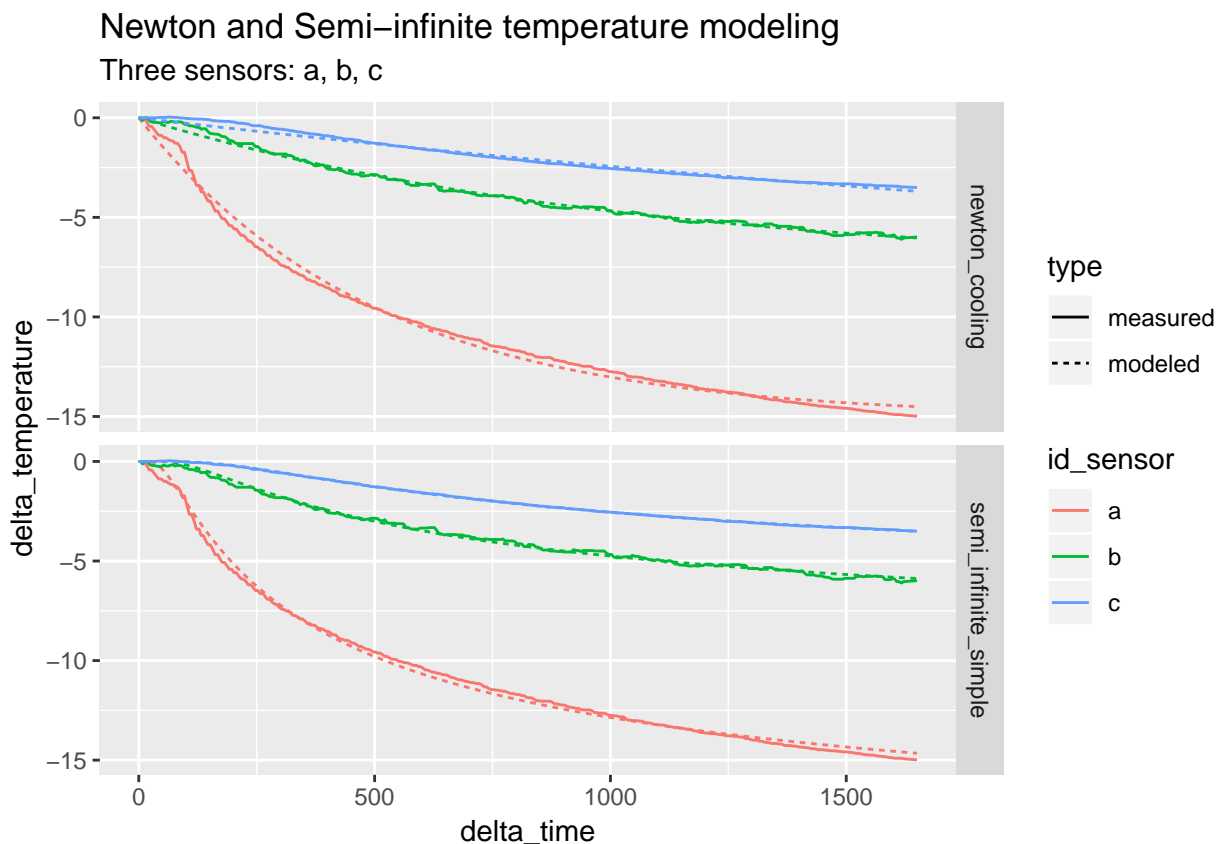
```

5.5.9 Visualize the predictions

```

predict_tall %>%
  ggplot(aes(x = delta_time, y = delta_temperature)) +
  geom_line(aes(color = id_sensor, linetype = type)) +
  facet_grid(id_model ~ .) +
  labs(title = "Newton and Semi-infinite temperature modeling",
        subtitle = "Three sensors: a, b, c")

```



5.5.10 Let's get the residuals

```

resid <-
  model_nested_new %>%
  mutate(resid = map(model, resid)) %>%

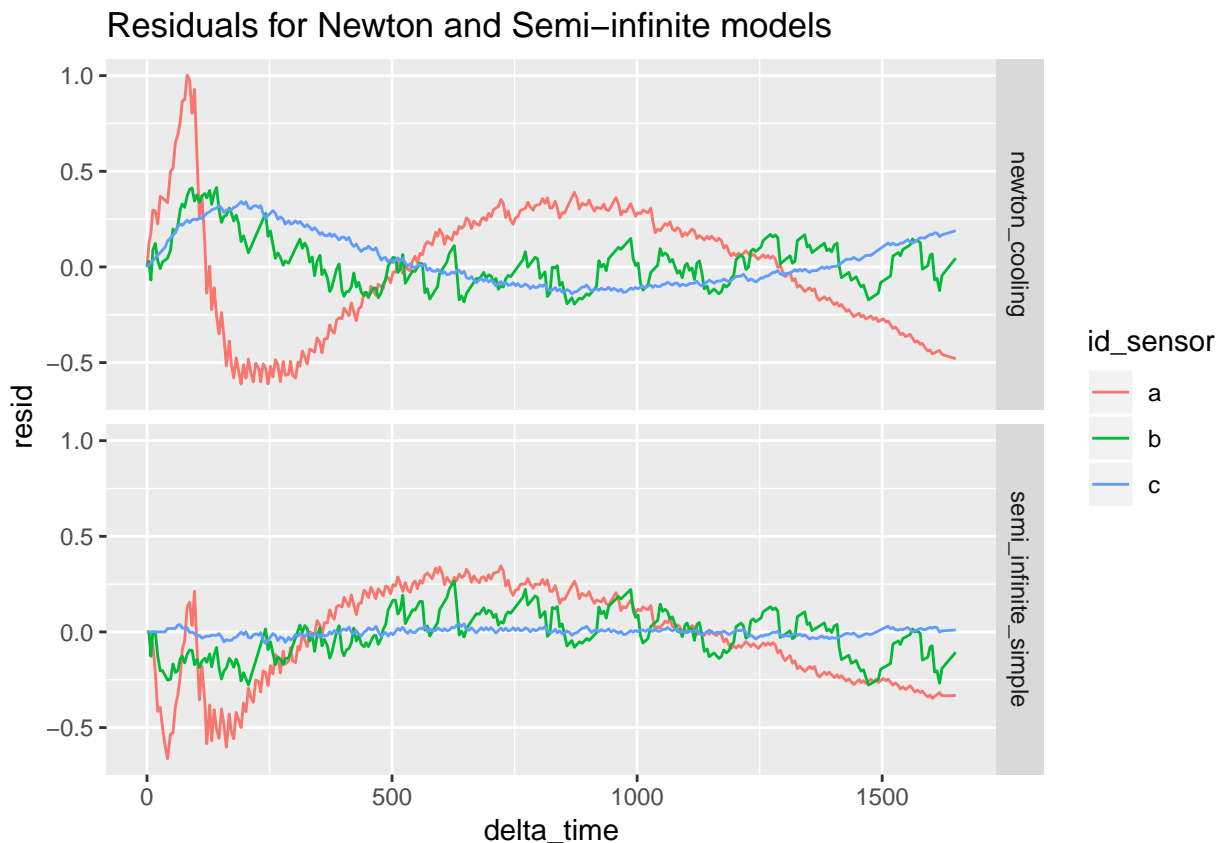
```

```
unnest(data, resid) %>%
print()
```

```
# A tibble: 1,962 x 5
  id_model      id_sensor resid delta_time delta_temperature
  <chr>        <chr>    <dbl>    <dbl>          <dbl>
1 newton_cooling a         0         0            0
2 newton_cooling a      0.120         4            0
3 newton_cooling a      0.179         8       -0.06
4 newton_cooling a      0.297        12       -0.06
5 newton_cooling a      0.292        17      -0.211
6 newton_cooling a      0.225        22      -0.423
7 newton_cooling a      0.369        27      -0.423
8 newton_cooling a      0.360        32      -0.574
9 newton_cooling a      0.348        37      -0.726
10 newton_cooling a      0.335        42      -0.878
# ... with 1,952 more rows
```

5.5.11 And visualize them

```
resid %>%
  ggplot(aes(x = delta_time, y = resid)) +
  geom_line(aes(color = id_sensor)) +
  facet_grid(id_model ~ .) +
  labs(title = "Residuals for Newton and Semi-infinite models")
```



5.6 Using broom package to look at model-statistics

We will use a previous defined dataframe with the model and data:

```
model_nested_new %>%
  print()

# A tibble: 6 x 4
  id_model      id_sensor data      model
  <chr>        <chr>    <list>    <list>
1 newton_cooling a      <tibble [327 x 2]> <nls>
2 newton_cooling b      <tibble [327 x 2]> <nls>
3 newton_cooling c      <tibble [327 x 2]> <nls>
4 semi_infinite_simple a    <tibble [327 x 2]> <nls>
5 semi_infinite_simple b    <tibble [327 x 2]> <nls>
6 semi_infinite_simple c    <tibble [327 x 2]> <nls>
```

The tidy() function extracts statistics from a model.

```
# apply over model_nested_new but only three variables
model_parameters <-
  model_nested_new %>%
  select(id_model, id_sensor, model) %>%
  mutate(tidy = map(model, tidy)) %>%
  select(-model) %>%
  unnest() %>%
  print()

# A tibble: 12 x 7
  id_model      id_sensor term      estimate std.error statistic  p.value
  <chr>        <chr>    <chr>      <dbl>    <dbl>    <dbl>    <dbl>
1 newton_cooli~ a      delta_te~ -15.1     0.0526   -286.    0.
2 newton_cooli~ a      tau_0      500.      4.84     103.    1.07e-250
3 newton_cooli~ b      delta_te~ -7.59     0.0676   -112.    6.38e-262
4 newton_cooli~ b      tau_0     1041.     16.2     64.2    9.05e-187
5 newton_cooli~ c      delta_te~ -9.87     0.704    -14.0    3.16e- 35
6 newton_cooli~ c      tau_0     3525.     299.     11.8    5.61e- 27
7 semi_infini~ a      delta_te~ -21.5     0.0649   -332.    0.
8 semi_infini~ a      tau_0     139.      1.15     121.    2.14e-272
9 semi_infini~ b      delta_te~ -10.6     0.0515   -206.    0.
10 semi_infini~ b      tau_0     287.      2.58     111.    1.46e-260
11 semi_infini~ c      delta_te~ -8.04     0.0129   -626.    0.
12 semi_infini~ c      tau_0     500.      1.07     468.    0.
```

5.6.1 Get a sense of the coefficients

```
model_summary <-
  model_parameters %>%
  select(id_model, id_sensor, term, estimate) %>%
  spread(key = "term", value = "estimate") %>%
  print()

# A tibble: 6 x 4
  id_model      id_sensor delta_temperature_0 tau_0
  <chr>        <chr>          <dbl> <dbl>
1 newton_cooli~ a      delta_temperature_0 500.
2 newton_cooli~ a      tau_0                4.84
3 newton_cooli~ b      delta_temperature_0 -7.59
4 newton_cooli~ b      tau_0                16.2
5 newton_cooli~ c      delta_temperature_0 -9.87
6 newton_cooli~ c      tau_0                299.
```

```
1 newton_cooling      a      -15.1  500.
2 newton_cooling      b      -7.59 1041.
3 newton_cooling      c      -9.87 3525.
4 semi_infinite_simple a     -21.5  139.
5 semi_infinite_simple b     -10.6  287.
6 semi_infinite_simple c      -8.04  500.
```

5.6.2 Summary

- this is just a small part of purrr
- there seem to be parallels between `tidyr::nest()/purrr::map()` and `dplyr::group_by()/dplyr::do()`
 - to my mind, the purrr framework is more understandable
 - update tweet from Hadley

References from Hadley:

- purrr 0.1.0 announcement
- purrr 0.2.0 announcement
- chapter from Garrett Golemund and Hadley's forthcoming book

Chapter 6

Logistic Regression. Diabetes

6.1 Introduction

Source: <https://github.com/AntoineGuillot2/Logistic-Regression-R/blob/master/script.R> Source: <http://enhancedatascience.com/2017/04/26/r-basics-logistic-regression-with-r/> Data: <https://www.kaggle.com/uciml/pima-indians-diabetes-database>

The goal of logistic regression is to predict whether an outcome will be positive (aka 1) or negative (i.e: 0). Some real life example could be:

- Will Emmanuel Macron win the French Presidential election or will he lose?
- Does Mr.X has this illness or not?
- Will this visitor click on my link or not?

So, logistic regression can be used in a lot of binary classification cases and will often be run before more advanced methods. For this tutorial, we will use the diabetes detection dataset from Kaggle.

This dataset contains data from Pima Indians Women such as the number of pregnancies, the blood pressure, the skin thickness, ... the goal of the tutorial is to be able to detect diabetes using only these measures.

6.2 Exploring the data

As usual, first, let's take a look at our data. You can download the data here then please put the file diabetes.csv in your working directory. With the summary function, we can easily summarise the different variables:

```
library(ggplot2)
library(data.table)

DiabetesData <- data.table(read.csv(file.path(data_raw_dir, 'diabetes.csv')))
```

Quick data summary

```
summary(DiabetesData)
```

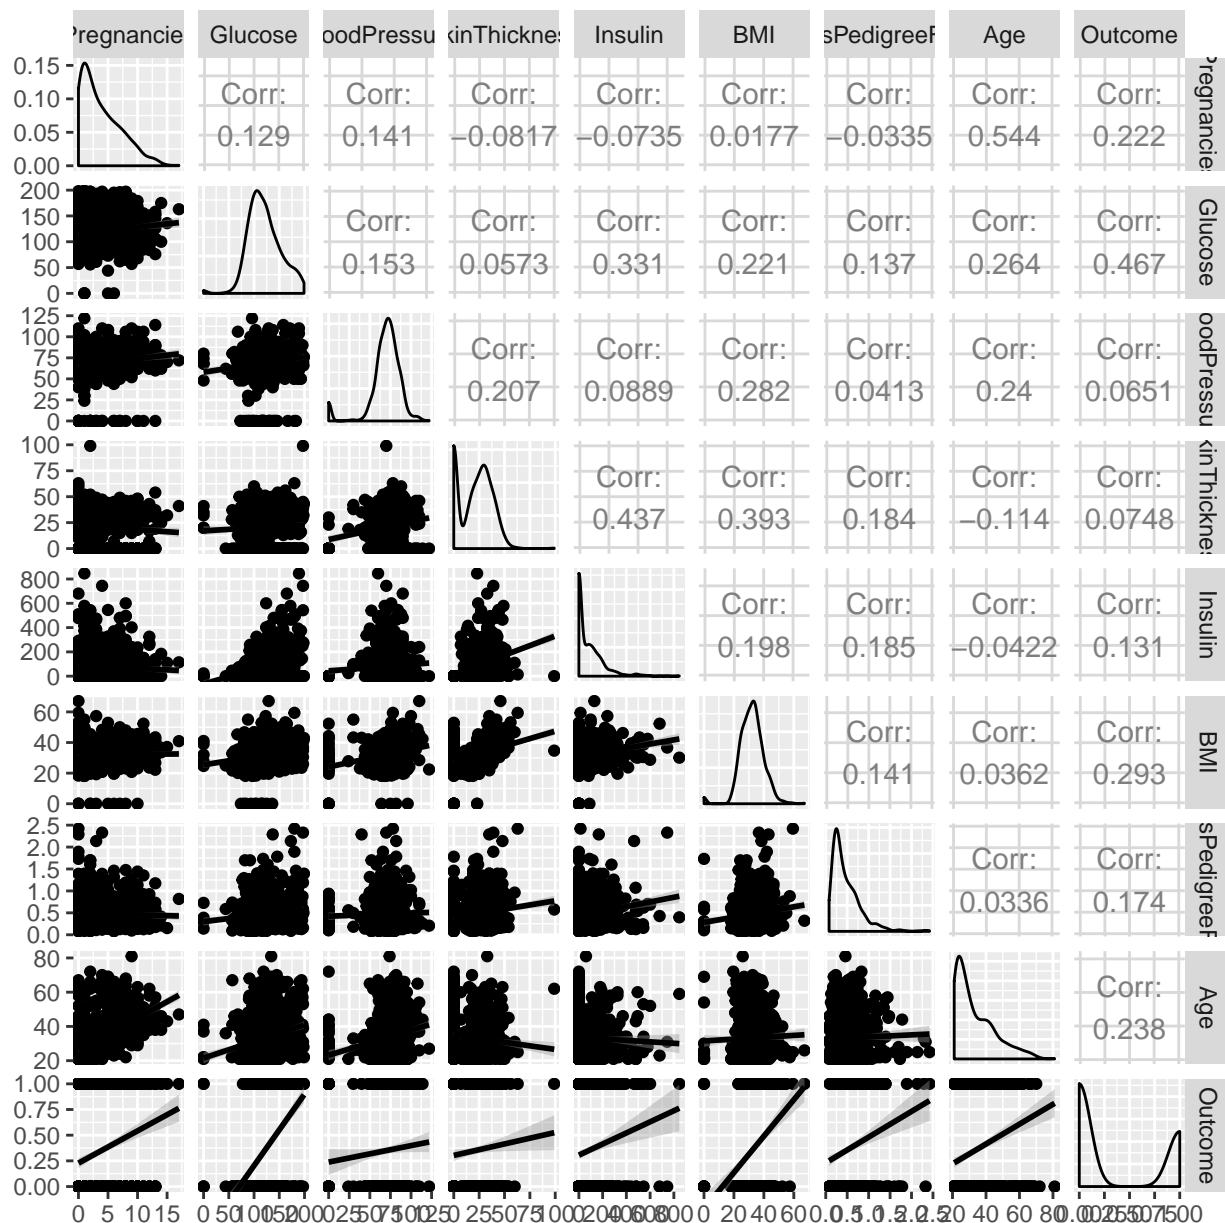
Pregnancies	Glucose	BloodPressure	SkinThickness
Min. : 0.000	Min. : 0.0	Min. : 0.00	Min. : 0.00
1st Qu.: 1.000	1st Qu.: 99.0	1st Qu.: 62.00	1st Qu.: 0.00
Median : 3.000	Median :117.0	Median : 72.00	Median :23.00
Mean : 3.845	Mean :120.9	Mean : 69.11	Mean :20.54

3rd Qu.: 6.000	3rd Qu.:140.2	3rd Qu.: 80.00	3rd Qu.:32.00
Max. :17.000	Max. :199.0	Max. :122.00	Max. :99.00
Insulin	BMI	DiabetesPedigreeFunction	Age
Min. : 0.0	Min. : 0.00	Min. :0.0780	Min. :21.00
1st Qu.: 0.0	1st Qu.:27.30	1st Qu.:0.2437	1st Qu.:24.00
Median : 30.5	Median :32.00	Median :0.3725	Median :29.00
Mean : 79.8	Mean :31.99	Mean :0.4719	Mean :33.24
3rd Qu.:127.2	3rd Qu.:36.60	3rd Qu.:0.6262	3rd Qu.:41.00
Max. :846.0	Max. :67.10	Max. :2.4200	Max. :81.00
Outcome			
Min. :0.000			
1st Qu.:0.000			
Median :0.000			
Mean :0.349			
3rd Qu.:1.000			
Max. :1.000			

The mean of the outcome is 0.35 which shows an imbalance between the classes. However, the imbalance should not be too strong to be a problem.

To understand the relationship between variables, a Scatter Plot Matrix will be used. To plot it, the GGally package was used.

```
# Scatter plot matrix
library(GGally)
ggpairs(DiabetesData, lower = list(continuous='smooth'))
```

The correlations between explanatory variables do not seem too strong. Hence the model is not likely to suffer from multicollinearity. All explanatory variable are correlated with the Outcome. At first sight, glucose rate is the most important factor to detect the outcome.

6.3 Logistic regression with R

After variable exploration, a first model can be fitted using the `glm` function. With `stargazer`, it is easy to get nice output in ASCII or even LaTeX.

```
# first model: all features
glm1 = glm(Outcome~.,
           DiabetesData,
           family = binomial(link="logit"))
```

```
summary(glm1)
```

Call:

```
glm(formula = Outcome ~ ., family = binomial(link = "logit"),
    data = DiabetesData)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.5566	-0.7274	-0.4159	0.7267	2.9297

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-8.4046964	0.7166359	-11.728	< 2e-16 ***
Pregnancies	0.1231823	0.0320776	3.840	0.000123 ***
Glucose	0.0351637	0.0037087	9.481	< 2e-16 ***
BloodPressure	-0.0132955	0.0052336	-2.540	0.011072 *
SkinThickness	0.0006190	0.0068994	0.090	0.928515
Insulin	-0.0011917	0.0009012	-1.322	0.186065
BMI	0.0897010	0.0150876	5.945	2.76e-09 ***
DiabetesPedigreeFunction	0.9451797	0.2991475	3.160	0.001580 **
Age	0.0148690	0.0093348	1.593	0.111192

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 993.48 on 767 degrees of freedom
 Residual deviance: 723.45 on 759 degrees of freedom
 AIC: 741.45

Number of Fisher Scoring iterations: 5

```
require(stargazer)
stargazer(glm1, type='text')
```

```
=====
                        Dependent variable:
                        -----
                                Outcome
                        -----
Pregnancies                0.123***
                             (0.032)

Glucose                    0.035***
                             (0.004)

BloodPressure              -0.013**
                             (0.005)

SkinThickness              0.001
                             (0.007)
```

Insulin	-0.001 (0.001)
BMI	0.090*** (0.015)
DiabetesPedigreeFunction	0.945*** (0.299)
Age	0.015 (0.009)
Constant	-8.405*** (0.717)

```
-----
Observations          768
Log Likelihood        -361.723
Akaike Inf. Crit.     741.445
=====
```

Note: *p<0.1; **p<0.05; ***p<0.01

The overall model is significant. As expected the glucose rate has the lowest p-value of all the variables. However, Age, Insulin and Skin Thickness are not good predictors of Diabetes.

6.4 A second model

Since some variables are not significant, removing them is a good way to improve model robustness. In the second model, SkinThickness, Insulin, and Age are removed.

```
# second model: selected features
glm2 = glm(Outcome~.,
           data = DiabetesData[,c(1:3,6:7,9), with=F],
           family = binomial(link="logit"))

summary(glm2)
```

Call:

```
glm(formula = Outcome ~ ., family = binomial(link = "logit"),
    data = DiabetesData[, c(1:3, 6:7, 9), with = F])
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.7931	-0.7362	-0.4188	0.7251	2.9555

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-7.954952	0.675823	-11.771	< 2e-16	***
Pregnancies	0.153492	0.027835	5.514	3.5e-08	***
Glucose	0.034658	0.003394	10.213	< 2e-16	***
BloodPressure	-0.012007	0.005031	-2.387	0.01700	*
BMI	0.084832	0.014125	6.006	1.9e-09	***
DiabetesPedigreeFunction	0.910628	0.294027	3.097	0.00195	**

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 993.48  on 767  degrees of freedom
Residual deviance: 728.56  on 762  degrees of freedom
AIC: 740.56

Number of Fisher Scoring iterations: 5
```

6.5 Classification rate and confusion matrix

Now that we have our model, let's access its performance.

```
# Correctly classified observations
mean((glm2$fitted.values>0.5)==DiabetesData$Outcome)
```

```
[1] 0.7747396
```

Around 77.4% of all observations are correctly classified. Due to class imbalance, we need to go further with a confusion matrix.

```
###Confusion matrix count
RP=sum((glm2$fitted.values>=0.5)==DiabetesData$Outcome & DiabetesData$Outcome==1)
FP=sum((glm2$fitted.values>=0.5)!=DiabetesData$Outcome & DiabetesData$Outcome==0)
RN=sum((glm2$fitted.values>=0.5)==DiabetesData$Outcome & DiabetesData$Outcome==0)
FN=sum((glm2$fitted.values>=0.5)!=DiabetesData$Outcome & DiabetesData$Outcome==1)
confMat<-matrix(c(RP,FP,FN,RN),ncol = 2)
colnames(confMat)<-c("Pred Diabetes", "Pred no diabetes")
rownames(confMat)<-c("Real Diabetes", "Real no diabetes")
confMat
```

	Pred Diabetes	Pred no diabetes
Real Diabetes	154	114
Real no diabetes	59	441

The model is good to detect people who do not have diabetes. However, its performance on ill people is not great (only 154 out of 268 have been correctly classified).

You can also get the percentage of Real/False Positive/Negative:

```
# Confusion matrix proportion
RPR=RP/sum(DiabetesData$Outcome==1)*100
FNR=FN/sum(DiabetesData$Outcome==1)*100
FPR=FP/sum(DiabetesData$Outcome==0)*100
RNR=RN/sum(DiabetesData$Outcome==0)*100
confMat<-matrix(c(RPR,FPR,FNR,RNR),ncol = 2)
colnames(confMat)<-c("Pred Diabetes", "Pred no diabetes")
rownames(confMat)<-c("Real Diabetes", "Real no diabetes")
confMat
```

	Pred Diabetes	Pred no diabetes
Real Diabetes	57.46269	42.53731
Real no diabetes	11.80000	88.20000

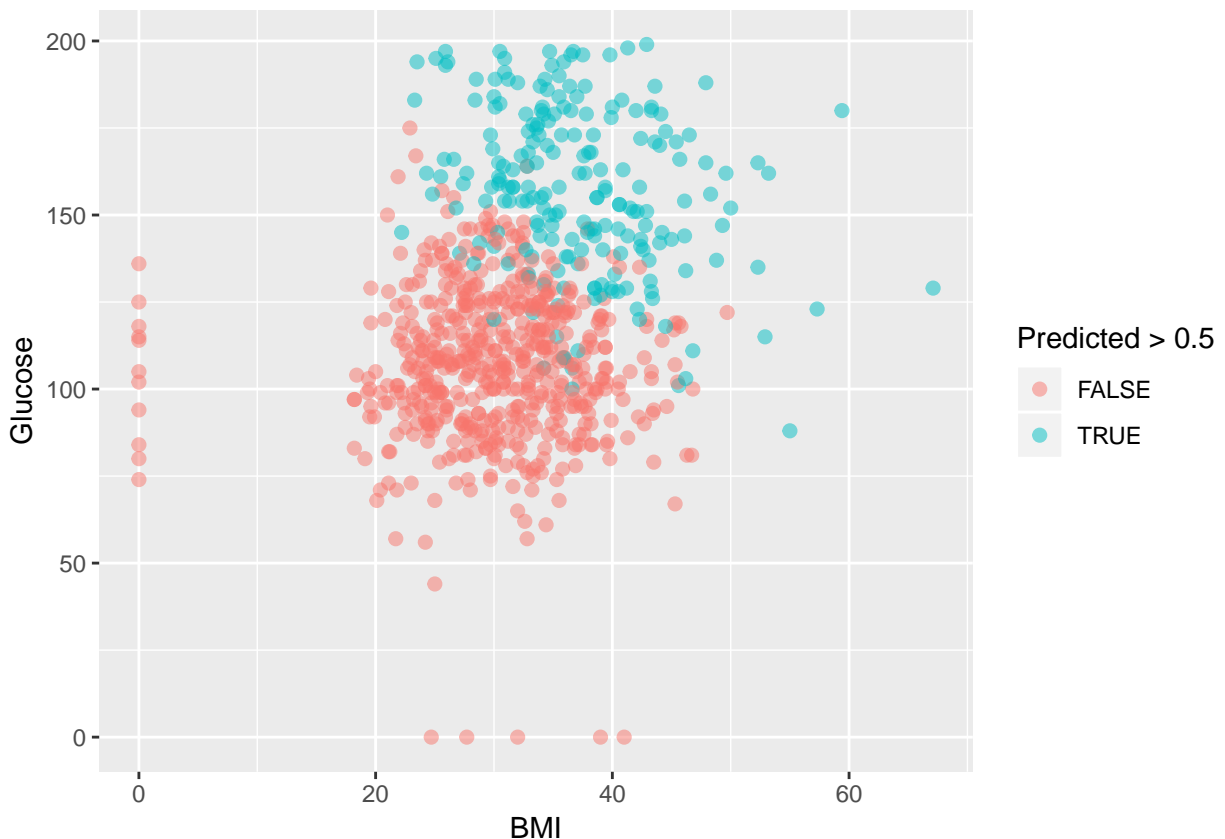
And here is the matrix, 57.5% of people with diabetes are correctly classified. A way to improve the false negative rate would lower the detection threshold. However, as a consequence, the false positive rate would increase.

6.6 Plots and decision boundaries

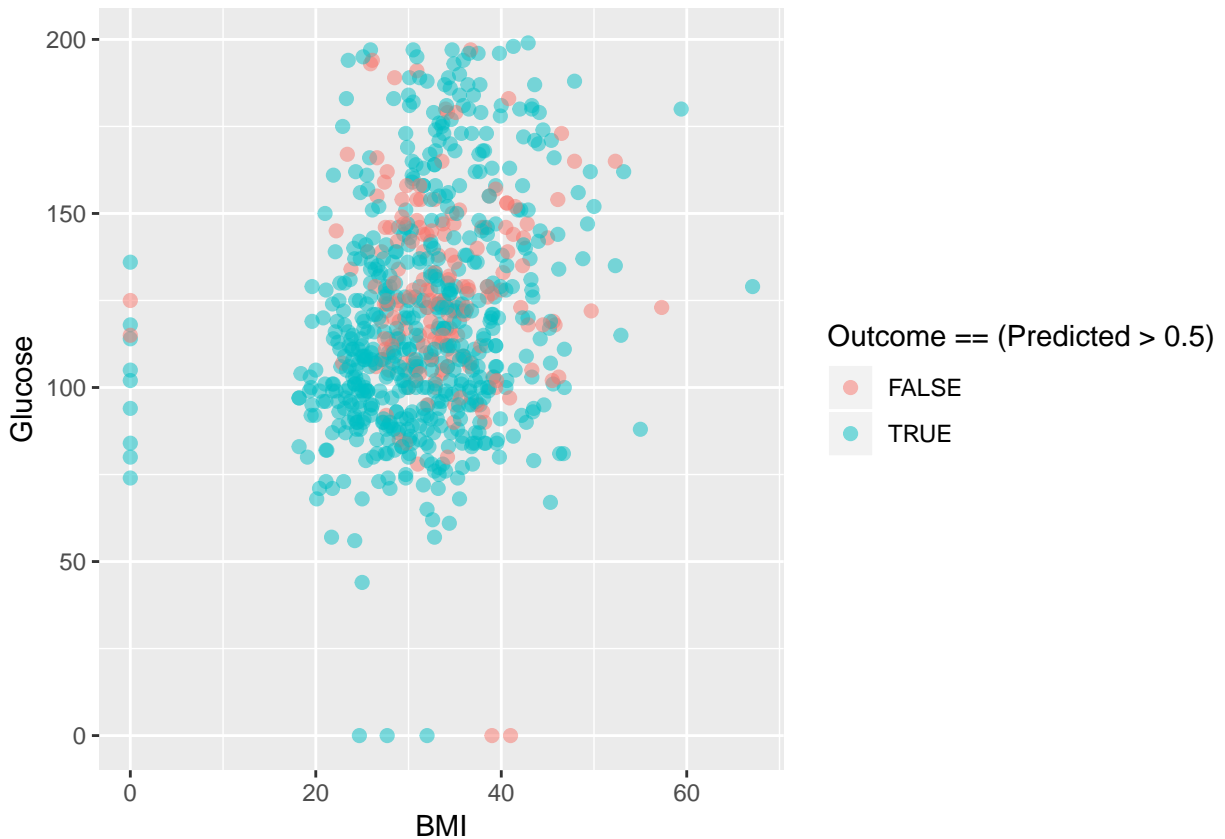
The two strongest predictors of the outcome are Glucose rate and BMI. High glucose rate and high BMI are strong indicators of Diabetes.

```
# Plot and decision boundaries
DiabetesData$Predicted <- glm2$fitted.values

ggplot(DiabetesData, aes(x = BMI, y = Glucose, color = Predicted > 0.5)) +
  geom_point(size=2, alpha=0.5)
```



```
ggplot(DiabetesData, aes(x=BMI, y = Glucose, color=Outcome == (Predicted > 0.5))) +
  geom_point(size=2, alpha=0.5)
```



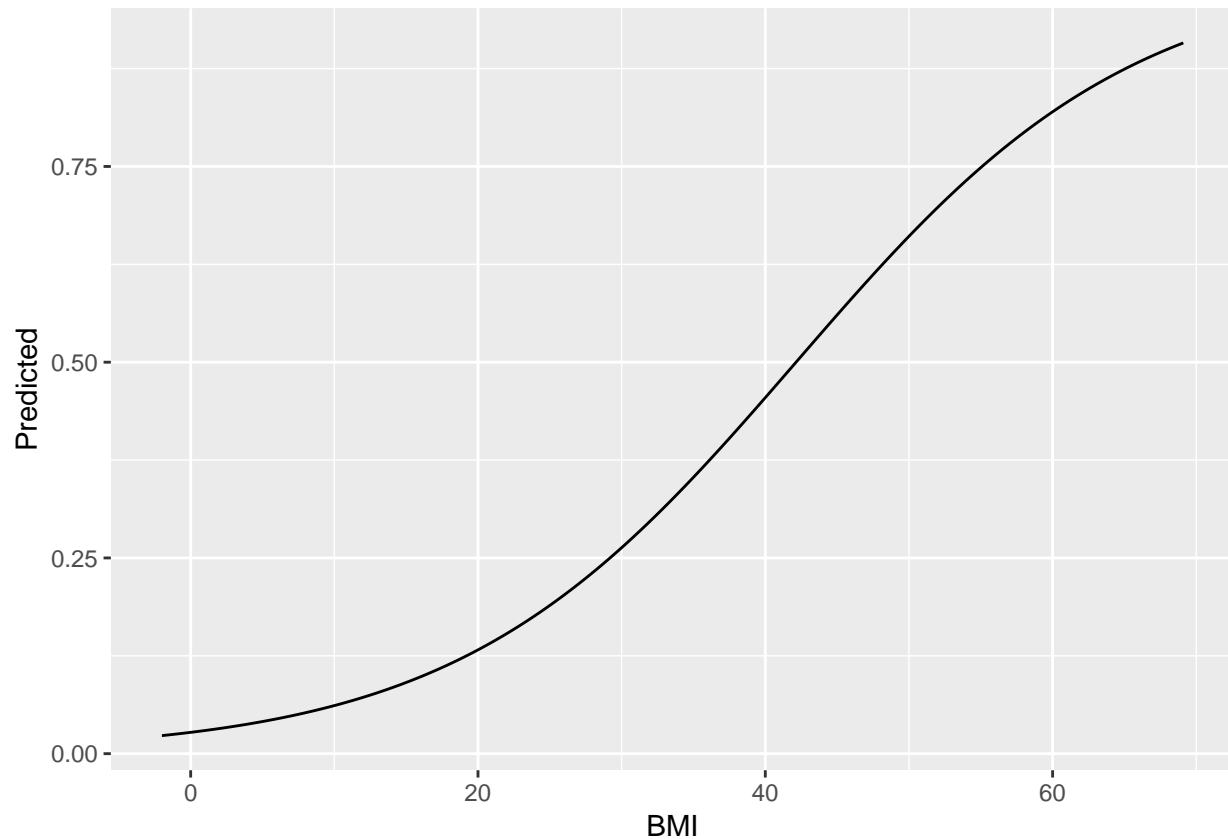
We can also plot both BMI and glucose against the outcomes, the other variables are taken at their mean level.

```
range(DiabetesData$BMI)
```

```
[1] 0.0 67.1
```

```
# BMI vs predicted
BMI_plot = data.frame(BMI = ((min(DiabetesData$BMI-2)*100):
                          (max(DiabetesData$BMI+2)*100))/100,
                      Glucose = mean(DiabetesData$Glucose),
                      Pregnancies = mean(DiabetesData$Pregnancies),
                      BloodPressure = mean(DiabetesData$BloodPressure),
                      DiabetesPedigreeFunction = mean(DiabetesData$DiabetesPedigreeFunction))

BMI_plot$Predicted = predict(glm2, newdata = BMI_plot, type = 'response')
ggplot(BMI_plot, aes(x = BMI, y = Predicted)) +
  geom_line()
```



```
range(BMI_plot$BMI)
```

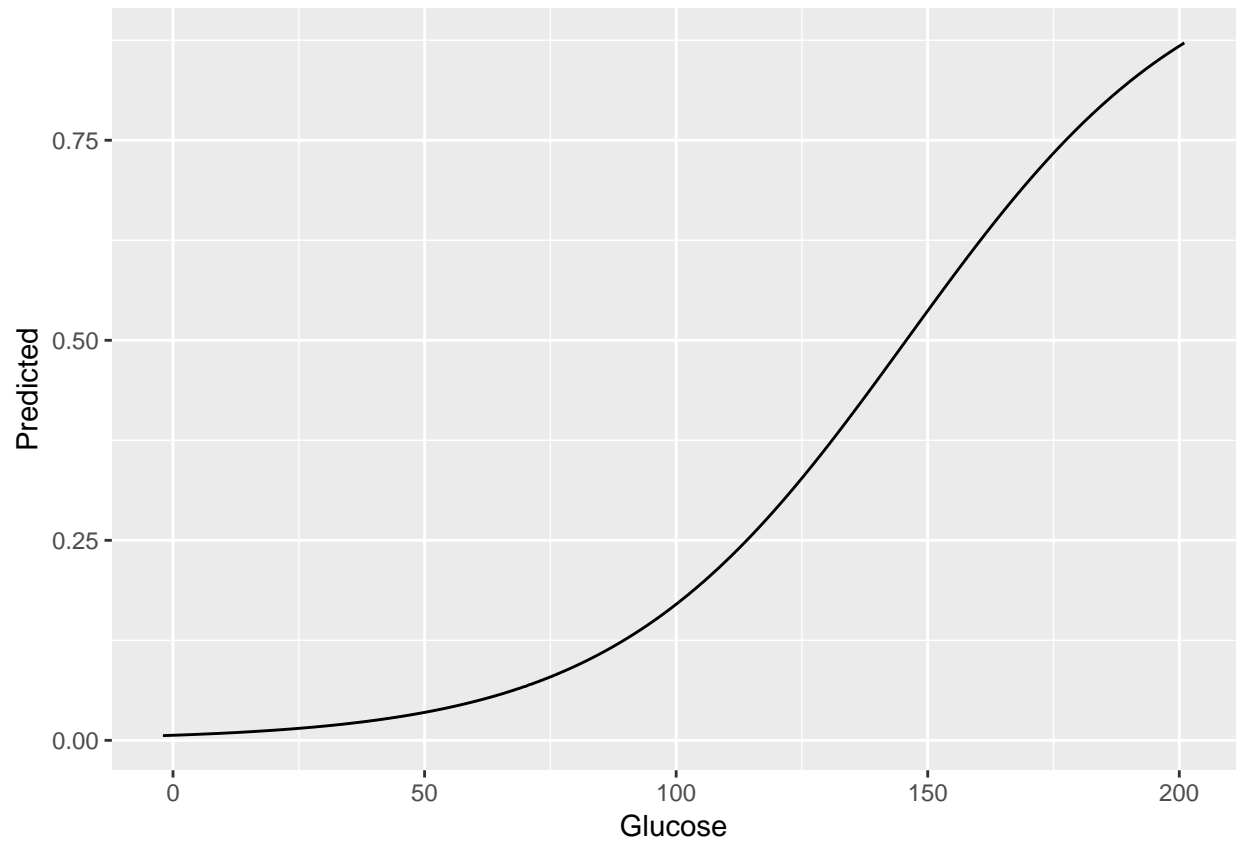
```
[1] -2.0 69.1
```

```
range(DiabetesData$Glucose)
```

```
[1] 0 199
```

```
# Glucose vs predicted
Glucose_plot=data.frame(Glucose =
  ((min(DiabetesData$Glucose-2)*100):
    (max(DiabetesData$Glucose+2)*100))/100,
  BMI=mean(DiabetesData$BMI),
  Pregnancies=mean(DiabetesData$Pregnancies),
  BloodPressure=mean(DiabetesData$BloodPressure),
  DiabetesPedigreeFunction=mean(DiabetesData$DiabetesPedigreeFunction))

Glucose_plot$Predicted = predict(glm2, newdata = Glucose_plot, type = 'response')
ggplot(Glucose_plot, aes(x = Glucose, y = Predicted)) +
  geom_line()
```



```
range(Glucose_plot$Glucose)
```

```
[1] -2 201
```